

1. Java Program to create TCP Sockets for C/S Communication

Server.java

```
import java.io.*;
import java.net.*;
public class MyServer {
    public static void main(String[] args){
        try{
            ServerSocket ss=new ServerSocket(6666);
            Socket s=ss.accept();//establishes connection
            DataInputStream dis=new DataInputStream(s.getInputStream());
            String str=(String)dis.readUTF();
            System.out.println("message= "+str);
            ss.close();
        }catch(Exception e){System.out.println(e);}
    }
}
```

Client.java

```
import java.io.*;
import java.net.*;
public class MyClient {
    public static void main(String[] args) {
        try{
            Socket s=new Socket("localhost",6666);
            DataOutputStream dout=new DataOutputStream(s.getOutputStream());
            dout.writeUTF("Hello Server");
            dout.flush();
            dout.close();
            s.close();
        }catch(Exception e){System.out.println(e);}
    }
}
```

2. Java Program to create UDP Sockets for C/S Communication

Server.java

```
import java.io.IOException;
import java.net.DatagramPacket;
import java.net.DatagramSocket;
import java.net.InetAddress;
import java.net.SocketException;
public class udpBaseServer_2{
    public static void main(String[] args) throws IOException{
        DatagramSocket ds = new DatagramSocket(1234);
        byte[] receive = new byte[65535];
        DatagramPacket DpReceive = null;
        while (true){
            DpReceive = new DatagramPacket(receive, receive.length);
            ds.receive(DpReceive);
            System.out.println("Client:-" + data(receive));
            if (data(receive).toString().equals("bye")){
                System.out.println("Client sent bye. ...EXITING");
                break;
            }
            receive = new byte[65535];}}
    public static StringBuilder data(byte[] a){
        if (a == null) return null;
        StringBuilder ret = new StringBuilder();
        int i = 0;
        while (a[i] != 0){
            ret.append((char) a[i]);
            i++;
        }
        return ret;}}
```

Client.java

```
import java.io.IOException;
import java.net.DatagramPacket;
import java.net.DatagramSocket;
import java.net.InetAddress;
import java.util.Scanner;
public class udpBaseClient_2{
    public static void main (String args[]) throws IOException{
        Scanner sc = new Scanner (System.in);
        DatagramSocket ds = new DatagramSocket ();
        InetAddress ip = InetAddress.getLocalHost ();
        byte buf[] = null;
        while (true){
            String inp = sc.nextLine ();
            buf = inp.getBytes ();
            DatagramPacket DpSend = new DatagramPacket (buf, buf.length, ip, 1234);
            ds.send (DpSend);
            if (inp.equals ("bye"))
                break;
        }
    }
}
```

3. Design a Web service using Simple Object Access Protocol (SOAP)

WebServiceserver.java

```
package vce.webservices.server;

import javax.xml.ws.Endpoint;

public class WebServiceServer {

    /**
     * Starts a simple server to deploy the web service.
     */
    public static void main(String[] args) {
        String bindingURI = "http://localhost:9898/md5WebService";
        MD5WebService webService = new MD5WebService();
        Endpoint.publish(bindingURI, webService);
        System.out.println("Server started at: " + bindingURI);
    }
}
```

webserviceClient.java

```
package vce.webservices.client;

public class WebServiceClient {

    /**
     * Starts the web service client.
     */
    public static void main(String[] args) {
        MD5WebServiceService client = new MD5WebServiceService();
        MD5WebService md5WebService = client.getMD5WebServicePort();
        String hash = md5WebService.hashString("hyderabad");
        System.out.println("MD5 hash string: " + hash);
    }
}
```

Md5webservice.java

```
package vce.webservices.server;

import java.security.MessageDigest;
import java.security.NoSuchAlgorithmException;

import javax.jws.WebMethod;
import javax.jws.WebService;

@WebService
public class MD5WebService {
    @WebMethod
    public String hashString(String input) {
        try {
            MessageDigest msgDigest = MessageDigest.getInstance("MD5");
            byte[] inputBytes = input.getBytes();
            byte[] hashedBytes = msgDigest.digest(inputBytes);

            StringBuffer sb = new StringBuffer();
            for (int i = 0; i < hashedBytes.length; i++) {
                sb.append(Integer.toString((hashedBytes[i] & 0xff) + 0x100, 16)
                    .substring(1));
            }

            return sb.toString();
        } catch (NoSuchAlgorithmException ex) {
            ex.printStackTrace();
            return "";
        }
    }
}
```

4. Developing a Multi chat application using Java.

```
import java.io.*;
import java.util.*;
import java.net.*;

public class Server{
    static Vector<ClientHandler> ar = new Vector<>();
    static int i = 0;
    public static void main(String[] args) throws IOException{
        ServerSocket ss = new ServerSocket(1234);
        Socket s;
        while (true){
            s = ss.accept();
            System.out.println("New client request received : " + s);
            DataInputStream dis = new DataInputStream(s.getInputStream());
            DataOutputStream dos = new DataOutputStream(s.getOutputStream());
            System.out.println("Creating a new handler for this client...");
            ClientHandler mtch = new ClientHandler(s,"client " + i, dis, dos);
            Thread t = new Thread(mtch);
            System.out.println("Adding this client to active client list");
            ar.add(mtch);
            t.start();
            i++;
        }
    }
}

class ClientHandler implements Runnable{
    Scanner scn = new Scanner(System.in);
    private String name;
    final DataInputStream dis;
    final DataOutputStream dos;
    Socket s;
    boolean isloggedin;
    public ClientHandler(Socket s, String name, DataInputStream dis, DataOutputStream dos) {
        this.dis = dis;
        this.dos = dos;
        this.name = name;
        this.s = s;
        this.isloggedin=true;
    }
    public void run() {
        String received;
        while (true){
            try{
                received = dis.readUTF();
                System.out.println(received);
                if(received.equals("logout")){
                    this.isloggedin=false;
                    this.s.close();
                    break;
                }
            }
        }
    }
}
```

```

        StringTokenizer st = new StringTokenizer(received, "#");
        String MsgToSend = st.nextToken();
        String recipient = st.nextToken();
        for (ClientHandler mc : Server.ar){
            if (mc.name.equals(recipient) && mc.isloggedin==true){
                mc.dos.writeUTF(this.name+" : "+MsgToSend);
                break;
            }
        }
    } catch (IOException e) {
        e.printStackTrace();
    }
}
try{
    this.dis.close();
    this.dos.close();

}catch(IOException e){
    e.printStackTrace();
}
}
}
}

```

5. Map reduce and Hadoop word frequency count

WCDriver.java

```

import java.io.IOException;
import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;
public class WCDriver extends Configured implements Tool {
    public int run(String args[]) throws IOException{
        if (args.length < 2) return -1;
        JobConf conf = new JobConf(WCDriver.class);
        FileInputFormat.setInputPaths(conf, new Path(args[0]));
        FileOutputFormat.setOutputPath(conf, new Path(args[1]));
        conf.setMapperClass(WCMapper.class);
        conf.setReducerClass(WCReducer.class);
        conf.setMapOutputKeyClass(Text.class);
        conf.setMapOutputValueClass(IntWritable.class);
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.class);
        JobClient.runJob(conf);
        return 0;
    }
    public static void main(String args[]) throws Exception{
        int exitCode = ToolRunner.run(new WCDriver(), args);
        System.out.println(exitCode);
    }
}

```

WCMapper.java

```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reporter;

public class WCMapper extends MapReduceBase implements Mapper<LongWritable, Text, Text,
IntWritable> {
    public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable> output,
Reporter rep) throws IOException
    {
        String line = value.toString();
        for (String word : line.split(" ")) {
            if (word.length() > 0)
                output.collect(new Text(word), new IntWritable(1));
        }
    }
}
```

WCReducer.java

```
import java.io.IOException;
import java.util.Iterator;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;

public class WCReducer extends MapReduceBase implements Reducer<Text, IntWritable, Text,
IntWritable> {
    public void reduce(Text key, Iterator<IntWritable> value, OutputCollector<Text, IntWritable>
output, Reporter rep) throws IOException {
        int count = 0;
        while (value.hasNext()) {
            IntWritable i = value.next();
            count += i.get();
        }
        output.collect(key, new IntWritable(count));
    }
}
```

Source Code:

Client.java

```
package TwoPC;
import java.io.*;
import java.net.*;
import java.awt.*;
import java.awt.event.*;
import javax.swing.*;
import java.sql.*;

class DBConnector {
    public static Connection getDBConnection(String dsn) throws Exception {
        Class.forName("sun.jdbc.odbc.JdbcOdbcDriver");
        return DriverManager.getConnection("jdbc:odbc:"+dsn);
    }
}

public class Client extends JFrame implements ActionListener{
    JButton b1,b2,b4,b5;
    JPanel p1,p2;
    JTextField t1;
    JLabel l1;
    ServerSocket ss;
    Socket s;
    DataOutputStream output;
    DataInputStream input;
    Connection con;
    Statement stmt;
    String serverMessage="Prepared";
    int port = 8890;
    String groupIP = "228.5.6.200";
    Client(){
        b1=new JButton("Prepared");
        b2=new JButton("NotPrepared");
        b4=new JButton("Execute");
        b5=new JButton("Exit");
        t1=new JTextField("",35);
        l1=new JLabel("SQL");
        p1=new JPanel();
        p2=new JPanel();
        p1.setLayout(new FlowLayout());
        p1.add(l1);
        p1.add(t1);
        p2.add(b1);
        p2.add(b2);
        p2.add(b4);
        p2.add(b5);
        add(p1);
        add(p2,"South");
        setSize(600,300);
        setTitle("Two Phase Commit Protocol: Client");
    }
}
```

```

        b1.addActionListener(this);
        b2.addActionListener(this);
        b4.addActionListener(this);
        b5.addActionListener(this);
        setVisible(true);
        setDefaultCloseOperation(EXIT_ON_CLOSE);
        MulticastSocket ms =null;
        InetAddress group ;
        try {
s = new Socket("localhost",8088);
System.out.println("Client Connected");
output=new DataOutputStream(s.getOutputStream());
input=new DataInputStream(s.getInputStream());
con = DBConnector.getDBConnection("my2pcdsn");
stmt = con.createStatement();
con.setAutoCommit(false);
ms = new MulticastSocket(port);
group= InetAddress.getByName(groupIP);
ms.joinGroup(group);
byte[] buffer = new byte[1024];
output.writeUTF("NotPrepared");
while (true) {
    DatagramPacket serMsg= new DatagramPacket(buffer, buffer.length);
    ms.receive(serMsg);
    String commitMsg = new String (serMsg.getData()).trim();
    if (commitMsg.equals("commit")) {
        System.out.println("Received "+commitMsg);
        con.commit();
        t1.setText("Transactions Committed");
        System.out.println("Transactions Committed");
    }
}
        }
        catch (ConnectException ce)
        {
            ce.printStackTrace();
            System.exit(0);
        }
        catch (Exception e)
        {
            e.printStackTrace();
        }
    }

    public void actionPerformed(ActionEvent ae){
        try
        {
            String str=ae.getActionCommand();

```



```

        if(str.equals("Execute")){
            String query = t1.getText();
            stmt.executeUpdate(query);
            t1.setText("Query Executed (NotPrepared)");
            output.writeUTF("NotPrepared");

        }

        if(str.equals("Prepared")){
            output.writeUTF("Prepared");
            t1.setText(input.readUTF());
        }

        if(str.equals("NotPrepared")){
            output.writeUTF("NotPrepared");
            t1.setText("NotPrepared");
        }

        if(str.equals("Exit")){
            output.writeUTF("Prepared");
            stmt.close();
            con.close();
            System.exit(0);
        }
    }
    catch(Exception e){
        JLabel errorFields = new JLabel("<HTML><FONT COLOR =
        BLUE>" + e.getMessage() + "</FONT></HTML>");
        JOptionPane.showMessageDialog(null,errorFields);
        e.printStackTrace();
    }
}

public static void main(String args[]){
    Client c=new Client();
}
}

```

Server.java

```
package TwoPC;
```

```
import java.io.*;
import java.net.*;
```

```
public class Server {

    public static ServerSocket ss;

    public Server() {
    }

    public static void main(String args[]) throws Exception {

```

```

        ss = new ServerSocket(8088);
        System.out.println("Two Phase Commit Protocol: Server");
        new Clients(2);
        while (true) {
            System.out.println("Server waiting: ");
            Socket s = ss.accept();
            new Coordinator(s);
        }
    }
}

```

```

class Clients {
    static int n;
    static String[] status;

    Clients(int num) {
        n = num;
        status = new String[n];
        for (int j = 0; j < n; j++) {
            status[j] = new String("");
        }
    }
}

```

```

class Coordinator implements Runnable {

    public static int i = -1;
    int flag = 1;
    Socket s;
    Thread t;
    MulticastSocket ms = null;
    InetAddress group;
    DataInputStream input;
    DataOutputStream output;
    int port = 8890;
    String groupIP = "228.5.6.200";
    Coordinator(Socket c) {
        s = c;
        try {
            input = new DataInputStream(s.getInputStream());
            output = new DataOutputStream(s.getOutputStream());
            ms = new MulticastSocket(port);
            group = InetAddress.getByName(groupIP);
            ms.joinGroup(group);
        } catch (Exception e) {
            e.printStackTrace();
        }
        t = new Thread(this);
        t.start();
        i++;
    }
}

```

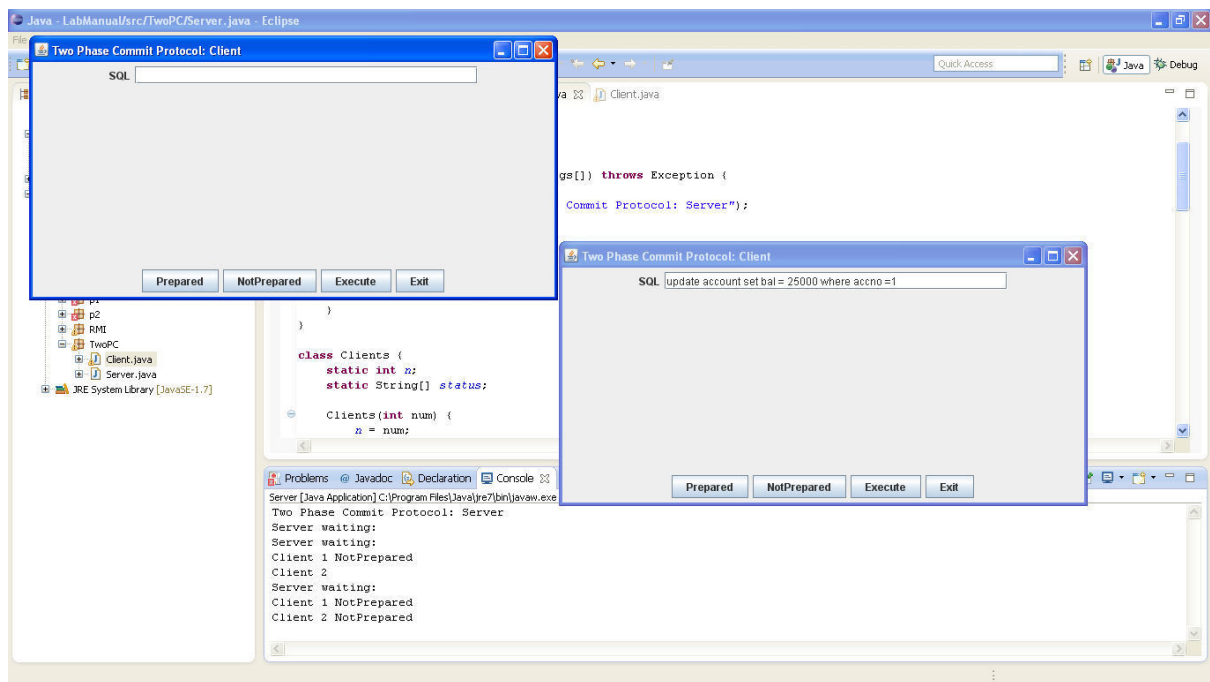
```

    }

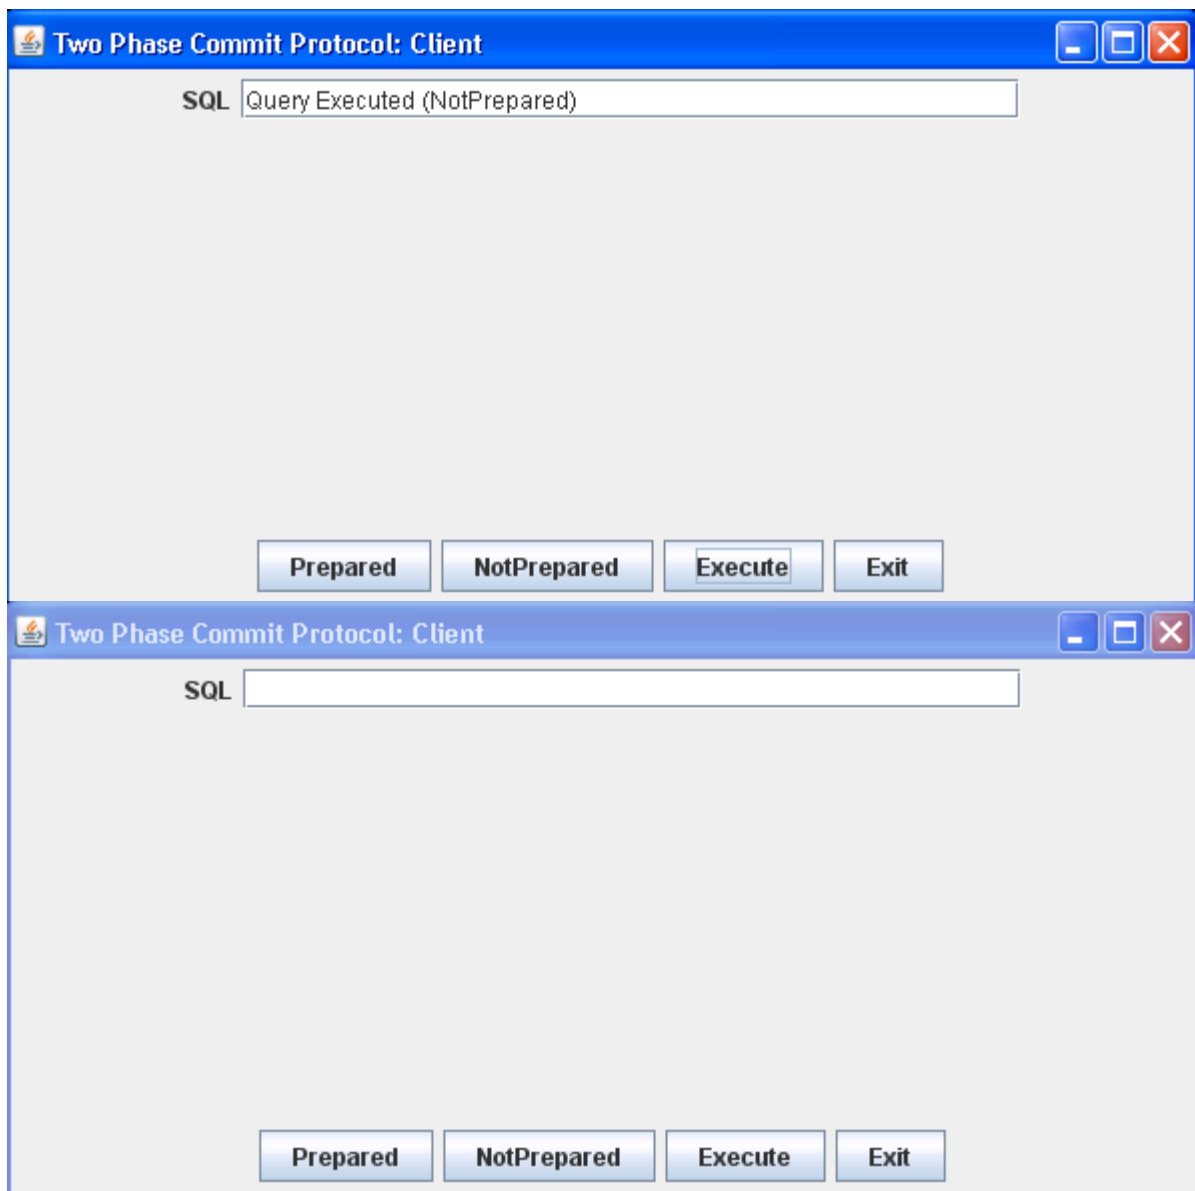
    public void run() {
        int index = i;
        String clientSattus;
        try {
            while (true) {
                clientSattus = input.readUTF();
                if (clientSattus.equalsIgnoreCase("Prepared")) {
                    output.writeUTF("Wait for others to prepare");
                }
                Clients.status[index] = new String(clientSattus);
                for (int k = 0; k < Clients.n; k++) {
                    System.out.println("Client " + (k + 1) + " " +
Clients.status[k]);
                    if (Clients.status[k].equalsIgnoreCase("Prepared"))
{
                        continue;
                    } else {
                        flag = 0;
                    }
                }
                if (flag == 1) {
                    byte[] msg = new String("commit").getBytes();
                    DatagramPacket msgpack = new
DatagramPacket(msg, msg.length, group, port);
                    ms.send(msgpack);
                    System.out.println("Commit message sent to clients:
" + new String(msg));
                }
                flag = 1;
            }
        } catch (Exception e) {
            e.printStackTrace();
        }
    }
}

```

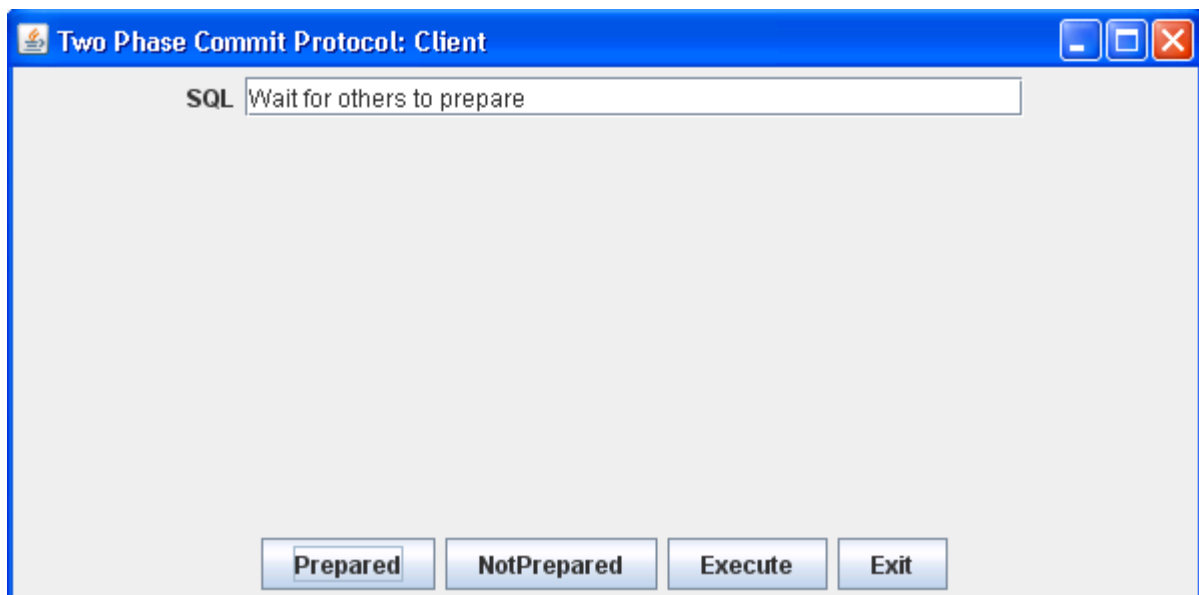
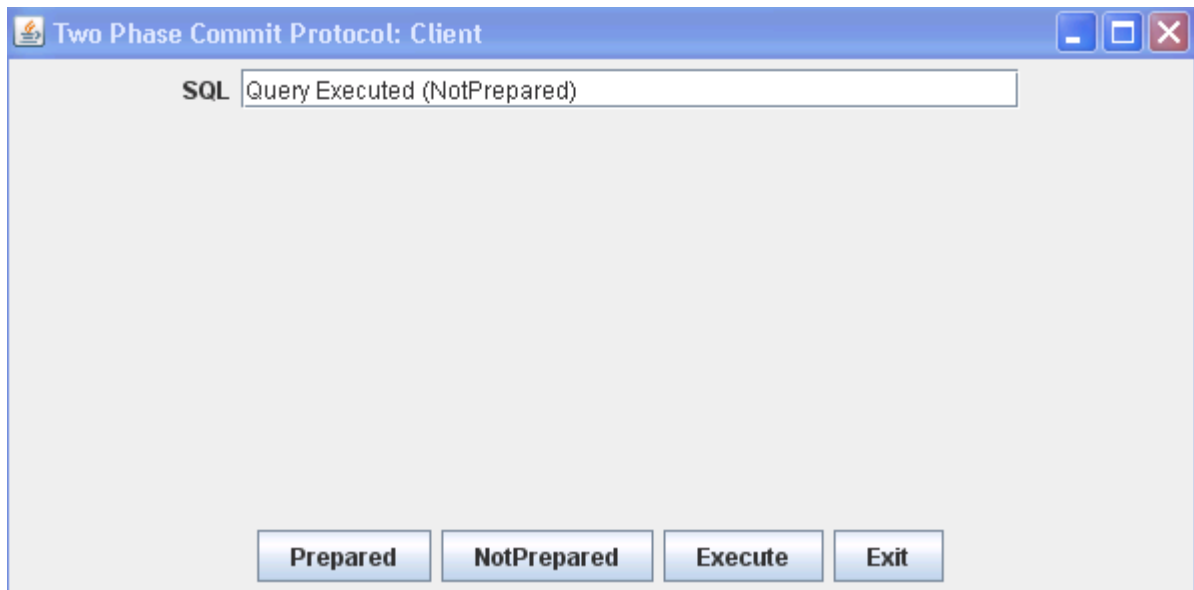
OUTPUT:



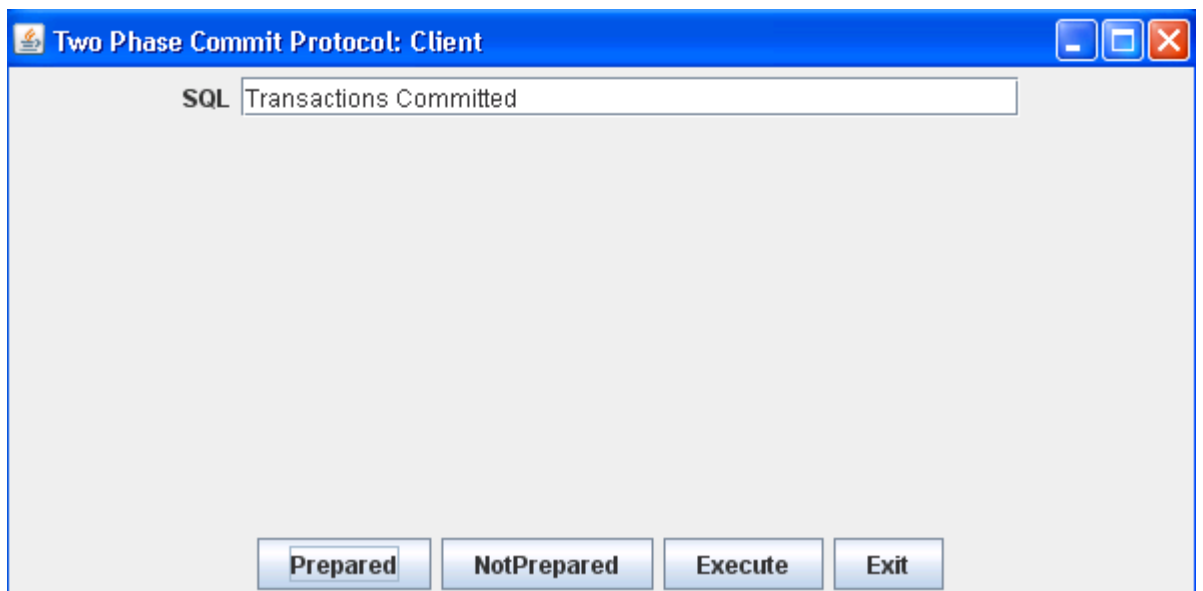
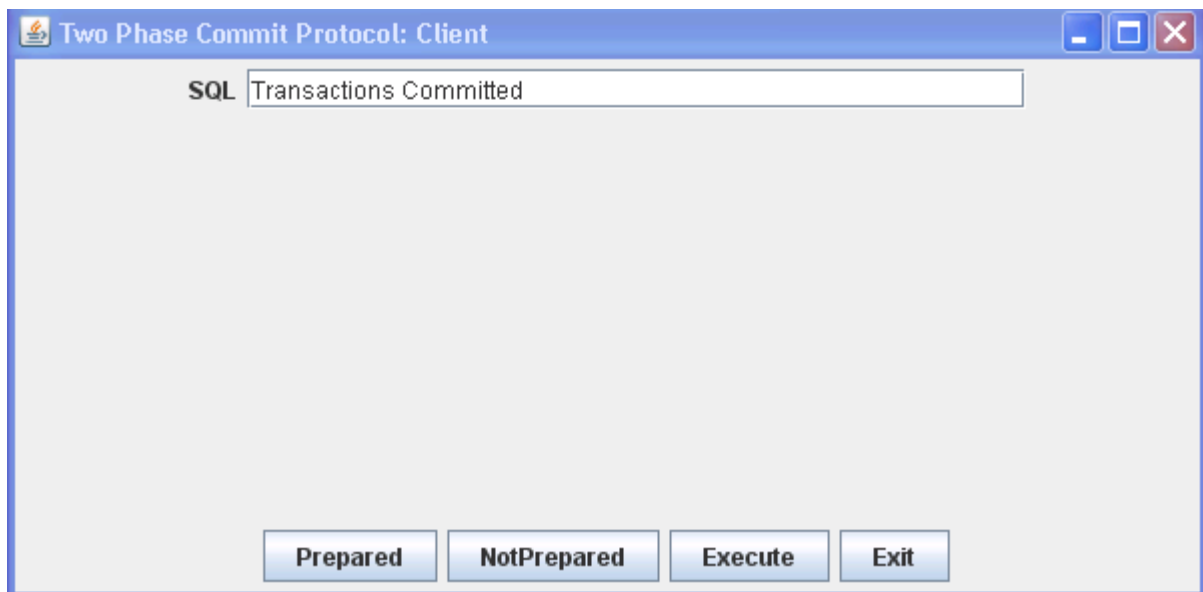
Client 1 executes Query

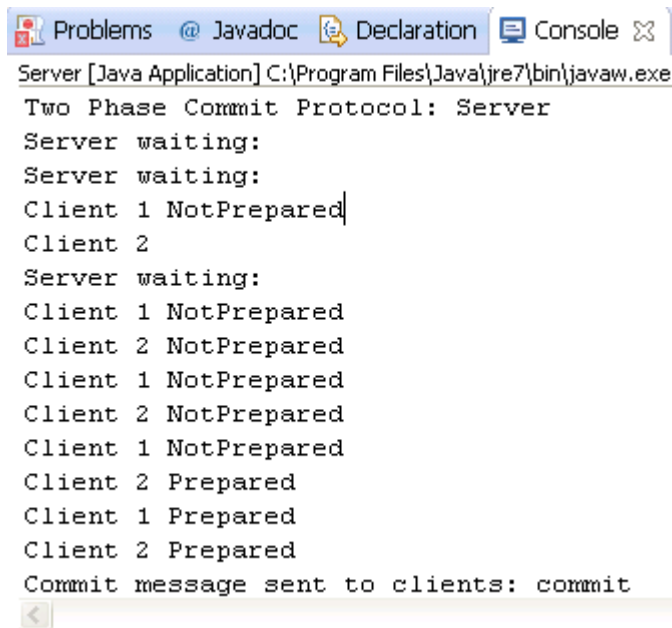


Client 2 says Prepared



Client 1 says Prepared and since both clients said Prepared, the transaction Commits



The image shows a screenshot of an IDE's console window. At the top, there is a tab bar with four tabs: 'Problems' (with a red bug icon), '@ Javadoc' (with a blue 'a' icon), 'Declaration' (with a yellow tag icon), and 'Console' (with a blue speech bubble icon). The 'Console' tab is selected. Below the tabs, the console output is displayed in a monospaced font. The text shows a sequence of messages related to a Two Phase Commit Protocol, including server waiting states, client preparedness reports, and a final commit message. A scrollbar is visible on the right side of the console area.

```
Server [Java Application] C:\Program Files\Java\jre7\bin\javaw.exe
Two Phase Commit Protocol: Server
Server waiting:
Server waiting:
Client 1 NotPrepared
Client 2
Server waiting:
Client 1 NotPrepared
Client 2 NotPrepared
Client 1 NotPrepared
Client 2 NotPrepared
Client 1 NotPrepared
Client 2 Prepared
Client 1 Prepared
Client 2 Prepared
Commit message sent to clients: commit
```


VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCC LAB

Name : _____ Roll No : 1602-19-733- Page No: _____

Lab Experiment

Hosting a Static Website

Accessing the AWS Management Console

1. At the top of these instructions, choose **Start Lab** to launch your lab.

A **Start Lab** panel opens, and it displays the lab status.

2. Wait until the **Start Lab** panel displays the message *Lab status: ready*, then close the panel by choosing the **X**.
3. At the top of these instructions, choose **AWS**.

This action opens the AWS Management Console in a new browser tab. The system automatically logs you in.

4. Arrange the **AWS Management Console** tab so that it displays alongside these instructions. Ideally, you will have both browser tabs open at the same time so that you can follow the lab steps more easily.

Do not change the Region unless specifically instructed to do so.

Task 1: Creating a bucket in Amazon S3

In this task, you will create an S3 bucket and configure it for static website hosting.

In the **AWS Management Console**, on the **Services** menu, choose **S3**.

5. Choose **Create bucket**

An S3 bucket name is globally unique, and the namespace is shared by all AWS accounts. After you create a bucket, the name of that bucket cannot be used by another AWS account in any AWS Region unless you delete the bucket.

Thus, for this lab, you will use a bucket name that includes a random number, such as: *website-123*

6. For **Bucket name**, enter: `website-<123>` (replace *<123>* with a random number)

Public access to buckets is blocked by default. Because the files in your static website will need to be accessible through the internet, you must permit public access.

- Verify the **AWS Region** is set to **us-east-1** (if it is not, choose the us-east-1 Region)

7. In the **Object Ownership** section, select **ACLs enabled**, then verify **Bucket owner preferred** is selected.

VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCC LAB

Name : _____ Roll No : 1602-19-733- Page No: _____

8. Clear **Block all public access**, then select the box that states **I acknowledge that the current settings may result in this bucket and the objects within becoming public**.
9. Choose **Create bucket**.

You can use tags to add additional information to a bucket, such as a project code, cost centre, or owner.

10. Choose the name of your new bucket.
11. Choose the **Properties** tab.
12. Scroll to the **Tags** panel.
13. Choose **Edit** then **Add tag** and enter:

- **Key:** Department
- **Value:** Marketing

14. Choose **Save changes** to save the tag.

Next, you will configure the bucket for static website hosting.

15. Stay in the **Properties** console.
16. Scroll to the **Static website hosting** panel.
17. Choose **Edit**

18. Configure the following settings:
 - **Static web hosting:** Enable
 - **Hosting type:** Host a static website
 - **Index document:** index.html
 - **Note:** You must enter this value, even though it is already displayed.
 - **Error document:** error.html

19. Choose **Save changes**
20. In the **Static website hosting** panel, choose the link under **Bucket website endpoint**.

You will receive a *403 Forbidden* message because the bucket permissions have not been configured yet. Keep this tab open in your web browser so that you can return to it later.

Your bucket has now been configured to host a static website.

Task 2: Uploading content to your bucket

In this task, you will upload the files that will serve as your static website to the bucket.

21. Right-click each of these links and download the files to your computer:

Ensure that each file keeps the same file name, including the extension.

- [index.html](#)

VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCC LAB

Name : _____ Roll No : 1602-19-733- Page No: _____

- [script.js](#)
- [style.css](#)

22. Return to the Amazon S3 console and in the `website-<123>` bucket you created earlier, choose the **Objects** tab.
23. Choose **Upload**.
24. Choose **Add files**
25. Locate and select the three files that you downloaded.
26. If prompted, choose I acknowledge that existing objects with the same name will be overwritten.
27. Choose **Upload**.

Your files are uploaded to the bucket.

- Choose **Close**

Task 3: Enabling access to the objects

Objects that are stored in Amazon S3 are private by default. This ensures that your organization's data remains secure.

In this task, you will make the uploaded objects publicly accessible.

First, confirm that the objects are currently private.

28. Return to the browser tab that showed the *403 Forbidden* message.
29. Refresh the webpage

You should still see a *403 Forbidden* message.

Analysis: This response is expected! This message indicates that your static website is being hosted by Amazon S3, but that the content is private.

You can make Amazon S3 objects public through two different ways:

- To make either a whole bucket public, or a specific directory in a bucket public, use a *bucket policy*.
 - To make individual objects in a bucket public, use an *access control list (ACL)*.
30. Return to the web browser tab with the Amazon S3 console (but do not close the website tab).
31. Select all three objects.
32. In the **Actions** menu, choose **Make public via ACL**.

A list of the three objects is displayed.

33. Choose **Make public**

VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCC LAB

Name : _____ Roll No : 1602-19-733- Page No: _____

Your static website is now publicly accessible.

34. Return to the web browser tab that has the *403 Forbidden* message.
35. Refresh the webpage.

You should now see the static website that is being hosted by Amazon S3.

Task 4: Updating the website

You can change the website by editing the HTML file and uploading it again to the S3 bucket.

36. On your computer, load the **index.html** file into a text editor (for example, Notepad or TextEdit).
37. Find the text **Served from Amazon S3** and replace it with **Created by <YOUR-NAME>**, substituting your name for **<YOUR-NAME>** (for example, *Created by Jane*).
38. Save the file.
39. Return to the Amazon S3 console and upload the **index.html** file that you just edited.
40. Select **index.html** and use the **Actions** menu to choose the **Make public via ACL** option again.
41. Return to the web browser tab with the static website and refresh the page.

Your name should now be on the page.

Your static website is now accessible on the internet. Because it is hosted on Amazon S3, the website has high availability and can serve high volumes of traffic without using any servers.

You can also use your own domain name to direct users to a static website that is hosted on Amazon S3. To accomplish this, you could use the Amazon Route 53 Domain Name System (DNS) service in combination with Amazon S3.

Submitting your work

42. At the top of these instructions, choose **Submit** to record your progress and when prompted, choose **Yes**.
43. If the results don't display after a couple of minutes, return to the top of these instructions, and choose **Grades**
44. To find detailed feedback on your work, choose **Details** followed by **View Submission Report**.

Lab complete

45. Choose **End Lab** at the top of this page, and then select **Yes** to confirm that you want to end the lab.

A panel indicates that *DELETE has been initiated... You may close this message box now*.

46. Select the **X** in the top right corner to close the panel.

VASAVI COLLEGE OF ENGINEERING

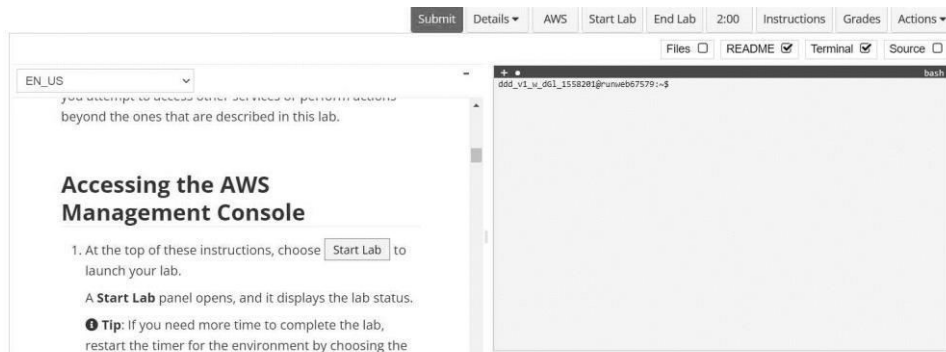
AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCC LAB

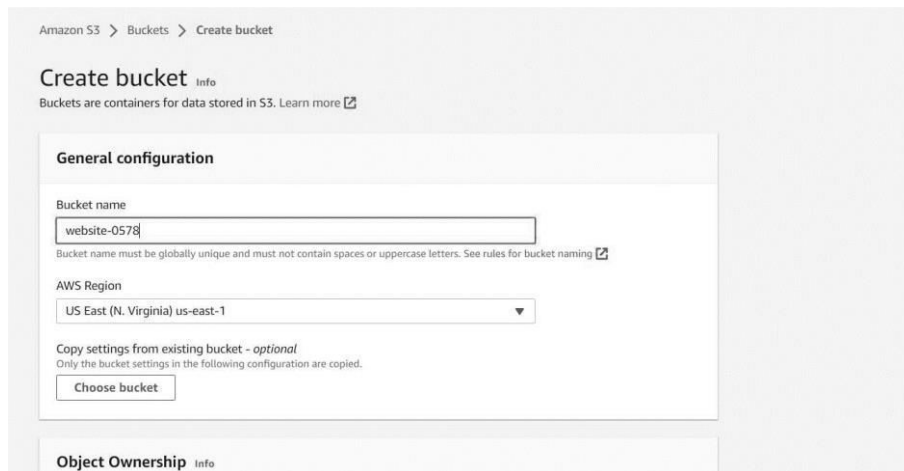
Name : _____ Roll No : 1602-19-733- Page No: _____

OUTPUT SCREENSHOTS:

Instructions Panel



Bucket Creation



Forbidden Error

403 Forbidden

- Code: AccessDenied
- Message: Access Denied
- RequestId: YJFWFVYPRMA8PRZ9
- HostId: ITJVNl0zuDIRbTQP0Uw9FgCeNEYkxaD4LJEnNSaNdCui0eRP4w4LqCrAnmZx5K8p0CIh+GwkMRQ=

An Error Occurred While Attempting to Retrieve a Custom Error Document

- Code: AccessDenied
- Message: Access Denied

VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCC LAB

Name : _____ Roll No : 1602-19-733- Page No: _____

Uploading Files

Upload succeeded
View details below.

Summary

Destination
s3://website-0578

Succeeded
3 Files, 14.8 KB (100.00%)

Failed
0 Files, 0 B (0%)

Files and folders

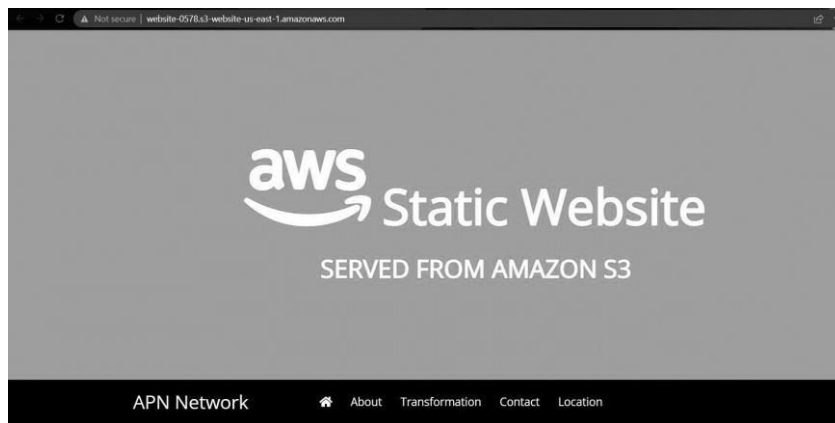
Configuration

Files and folders (3 Total, 14.8 KB)

Find by name

Name	Folder	Type	Size	Status	Error
index.html	-	text/html	9.8 KB	Succeeded	-
script.js	-	text/javascript	2.0 KB	Succeeded	-
style.css	-	text/css	3.1 KB	Succeeded	-

Static Website-1



Static Website-2 (Updated)



VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCC LAB

Name : _____ Roll No : 1602-19-733- Page No: _____

Lab Experiment

Introducing Amazon Elastic File System (Amazon EFS)

Accessing the AWS Management Console

1. At the top of these instructions, choose **Start Lab** to launch your lab.

A **Start Lab** panel opens, and it displays the lab status.

Tip: If you need more time to complete the lab, restart the timer for the environment by choosing the **Start Lab** button again.

2. Wait until the **Start Lab** panel displays the message *Lab status: ready*, then close the panel by choosing the **X**.
3. At the top of these instructions, choose **AWS**.
4. Arrange the **AWS Management Console** tab so that it displays alongside these instructions. Ideally, you will have both browser tabs open at the same time so that you can follow the lab steps more easily.

Task 1: Creating a security group to access your EFS file system

5. In the **AWS Management Console**, on the Services menu, choose **EC2**.
6. In the navigation pane on the left, choose **Security Groups**.
7. Copy the **Security group ID** of the *EFSCliant* security group to your text editor.

The Group ID should look similar to *sg-03727965651b6659b*.

8. Choose **Create security group** then configure:
 - **Security group name:** *EFS Mount Target*
 - **Description:** *Inbound NFS access from EFS clients*
 - **VPC:** *Lab VPC*
9. Under the **Inbound rules** section, choose **Add rule** then configure:
 - **Type:** *NFS*
 - **Source:**
 - *Custom*
 - In the *Custom* box, paste the security group's **Security group ID** that you copied to your text editor
 - Choose **Create security group**.

Task 2: Creating an EFS file system

10. On the Services menu, choose **EFS**.
11. Choose **Create file system**
12. In the **Create file system** window, choose **Customize**

VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCC LAB

Name : _____ Roll No : 1602-19-733- Page No: _____

13. On **Step 1**:

- Uncheck Enable automatic backups.
- **Lifecycle management**: Select *None*
- In the **Tags** section, configure:
 - **Key**: Name
 - **Value**: My First EFS File System

14. Choose Next

15. For **VPC**, select *Lab VPC*.

16. Detach the default security group from each *Availability Zone* mount target by choosing the check box on each default security group.

17. Attach the **EFS Mount Target** security group to each *Availability Zone* mount target by:

- Selecting each **Security groups** check box.
- Choosing **EFS Mount Target**

A mount target is created for each subnet

18. Choose Next

19. On **Step 3**, choose Next

20. On **Step 4**:

- Review your configuration.
- Choose **Create**

Proceed to the next step after the **Mount target state** for each mount target changes to *Available*. Choose the screen refresh button after 2–3 minutes to check its progress.

Task 3: Connecting to your EC2 instance via SSH

In this task, you will connect to your EC2 instance by using Secure Shell (SSH).

21. Above these instructions that you are currently reading, choose the **Details** dropdown menu, and then select **Show**

A **Credentials** window opens.

22. Choose the **Download PPK** button and save the **labsuser.ppk** file.

Note: Typically, your browser saves the file to the **Downloads** directory.

23. Note the **EC2PublicIP** address if it is displayed.

24. Exit the **Details** panel by choosing the **X**.

25. To use SSH to access the EC2 instance, you must use ***PuTTY***. If you do not have PuTTY installed on your computer, [download PuTTY](#).

26. Open **putty.exe**.

27. To keep the PuTTY session open for a longer period of time, configure the PuTTY timeout:

- Choose **Connection**
- **Seconds between keepalives**: 30

28. Configure your PuTTY session by using the following settings.

VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCC LAB

Name : _____ Roll No : 1602-19-733- Page No: _____

- Choose **Session**
- **Host Name (or IP address):** Paste the **EC2PublicIP** for the instance you noted earlier
 - Alternatively, return to the Amazon EC2 console and choose **Instances**
 - Select the instance you want to connect to
 - In the *Description* tab, copy the **IPv4 Public IP** value
- Back in PuTTY, in the **Connection** list, expand **SSH**
- Choose **Auth** (but don't expand it)
- Choose **Browse**
- Browse to the *labsuser.ppk* file that you downloaded, select it, and choose **Open**
- Choose **Open** again

29. To trust and connect to the host, choose **Yes**.

30. When you are prompted with **login as**, enter: `ec2-user`.

This action connects you to the EC2 instance.

Task 4: Creating a new directory and mounting the EFS file system

31. In your SSH session, make a new directory by entering `sudo mkdir efs`
32. Back in the **AWS Management Console**, on the Services menu, choose **EFS**.
33. Choose **My First EFS File System**.
34. In the **Amazon EFS Console**, on the top right corner of the page, choose **Attach** to open the Amazon EC2 mount instructions.
35. Copy the entire command in the **Using the NFS client** section.

The mount command should look similar to this example:

```
sudo mount -t nfs4 -o nfsvers=4.1,rsize=1048576,wsiz=1048576,hard,timeo=600,retrans=2,noresvport fs-bce57914.efs.us-west-2.amazonaws.com:/ efs
```

The provided `sudo mount...` command uses the default Linux mount options.

36. In your Linux SSH session, mount your Amazon EFS file system by:
 - Pasting the command
 - Pressing ENTER
37. Get a full summary of the available and used disk space usage by entering:

```
sudo df -hT
```

Task 5: Examining the performance behavior of your new EFS file system

38. Examine the write performance characteristics of your file system by entering:

VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCC LAB

Name : _____ Roll No : 1602-19-733- Page No: _____

```
sudo fio --name=fio-efs --filesize=10G --filename=./efs/fio-efs-test.img --bs=1M --nrfiles=1 --direct=1 -  
-sync=0 --rw=write --iodepth=200 --ioengine=libaio
```

Monitoring performance by using Amazon CloudWatch

39. In the **AWS Management Console**, on the Services menu, choose **CloudWatch**.
40. In the navigation pane on the left, choose **Metrics**.
41. In the **All-metrics** tab, choose **EFS**.
42. Choose **File System Metrics**.
43. Select the row that has the **PermittedThroughput** Metric Name.

You might need to wait 2–3 minutes and refresh the screen several times before all available metrics, including **PermittedThroughput**, calculate and populate.

44. On the graph, choose and drag around the data line. If you do not see the line graph, adjust the time range of the graph to display the period during which you ran the `fio` command.
45. Pause your pointer on the data line in the graph. The value should be *105M*.
46. In the **All-metrics** tab, *uncheck* the box for **PermittedThroughput**.
47. Select the check box for **DataWriteIOBytes**.

If you do not see *DataWriteIOBytes* in the list of metrics, use the **File System Metrics** search to find it.

48. Choose the **Graphed metrics** tab.
49. On the **Statistics** column, select **Sum**.
50. On the **Period** column, select **1 Minute**.
51. Pause your pointer on the peak of the line graph. Take this number (in bytes) and divide it by the duration in seconds (60 seconds). The result gives you the write throughput (B/s) of your file system during your test.

VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCC LAB

Name : _____ Roll No : 1602-19-733- Page No: _____

OUTPUT SCREENSHOTS:

Task-1: Creating a security group to access your EFS file system

The screenshot shows the 'Basic details' and 'Inbound rules' sections of the AWS IAM console for a new security group.

Basic details

- Security group name:
- Description:
- VPC:

Inbound rules

Type	Protocol	Port range	Source	Description - optional
NFS	TCP	2049	Custom <input type="text" value="sg-0ab07af222dc18773"/>	

Buttons: Add rule, Delete

Task-2: Creating an EFS file system

The screenshot shows the 'Virtual Private Cloud (VPC)' and 'Mount targets' sections of the AWS IAM console.

Virtual Private Cloud (VPC)

Choose the VPC where you want EC2 instances to connect to your file system. [Learn more](#)

VPC:

Mount targets

A mount target provides an NFSv4 endpoint at which you can mount an Amazon EFS file system. We recommend creating one mount target per Availability Zone. [Learn more](#)

Availability zone	Subnet ID	IP address	Security groups
us-east-1a	subnet-0099cdaaeac...	Automatic	<input type="text" value="sg-00f8ba5a438a74694"/>
us-east-1b	subnet-0ed63276a5...	Automatic	<input type="text" value="sg-00f8ba5a438a74694"/>

Buttons: Add mount target, Remove

Mount targets			
Availability zone	Subnet	IP address	Security groups
us-east-1a	subnet-0099cdaaeac24424e	-	sg-00f8ba5a438a74694
us-east-1b	subnet-0ed63276a5874d404	-	sg-00f8ba5a438a74694

VASAVI COLLEGE OF ENGINEERING

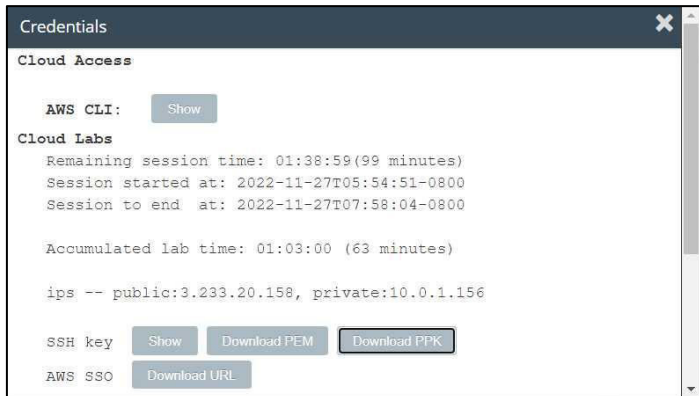
AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCC LAB

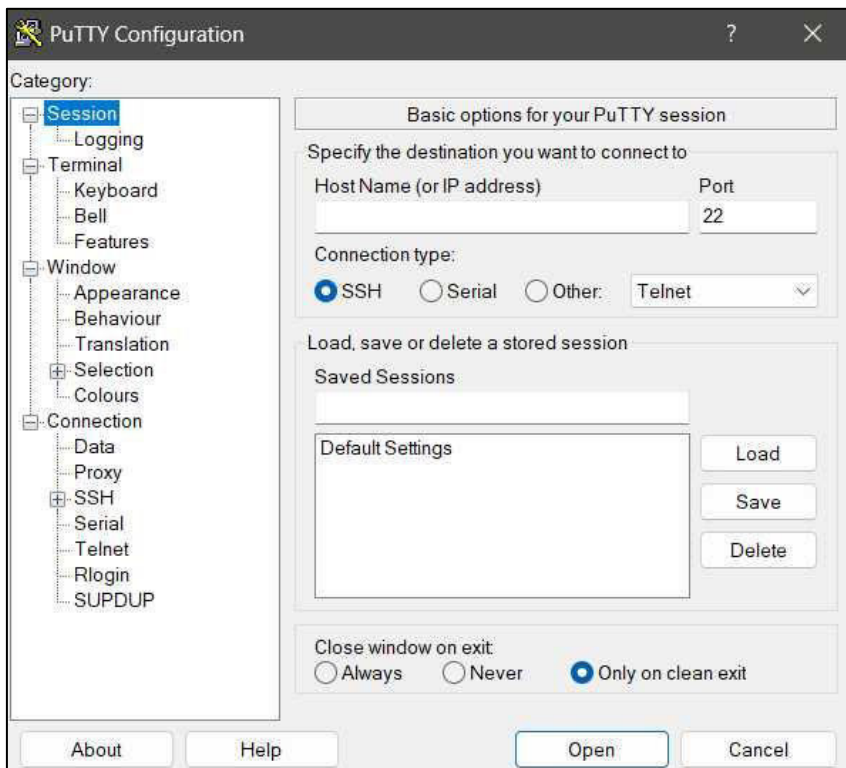
Name : _____ Roll No : 1602-19-733- Page No: ____

Task-3:

Credentials Tab



Putty Config



VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCC LAB

Name : _____ Roll No : 1602-19-733- Page No: _____

```
ec2-user@ip-10-0-1-156:~  
login as: ec2-user  
Authenticating with public key "imported-openssh-key"  
  
  _ | _ | _ )  
  _ | ( _ | _ /  Amazon Linux 2 AMI  
  _ | \ _ | _ |  
  
https://aws.amazon.com/amazon-linux-2/  
[ec2-user@ip-10-0-1-156 ~]$
```

Task-5: Examining the performance behavior of your new EFS file system



VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering

NAME OF THE LABORATORY : DSCCLAB

Name : _____ Roll No : 1602-19-733-0 Page No: _____

LAB PROGRAM

Experiment: Deploying a Node.js Web Application on AWS

HARDWARE REQUIREMENTS: Core I5 Processor, 4 GB RAM, 40GB HDD

SOFTWARE REQUIREMENTS: Amazon AWS, EC2, VS Code/Eclipse, Node, NPM, GIT, Putty

Description:

Node.js is a JavaScript runtime environment that allows one to run JS on the server. It is built on the open-source V8 JavaScript engine used in Chrome and written in C++ which executes JS in a standalone environment.

In this experiment, we clone a Nodejs application from GITHUB and deploy this application on to Amazon EC2 instance, make it available over Amazon AWS URI.

Steps to configure EC2 Instance :

1. Create an EC2 instance and Launch it:

Choose amazon Ec2 instance machine image as Ubuntu 18.04 64 bit with type of micro.

(Login to AwsAcademy,

LMS-Dashboard - AWS Academy Learner Lab – Educator

Click on Modules

Click on Learner Lab

Click on Start Lab

Click on AWS

Services – EC2

EC2 – Instances – Launch an instance

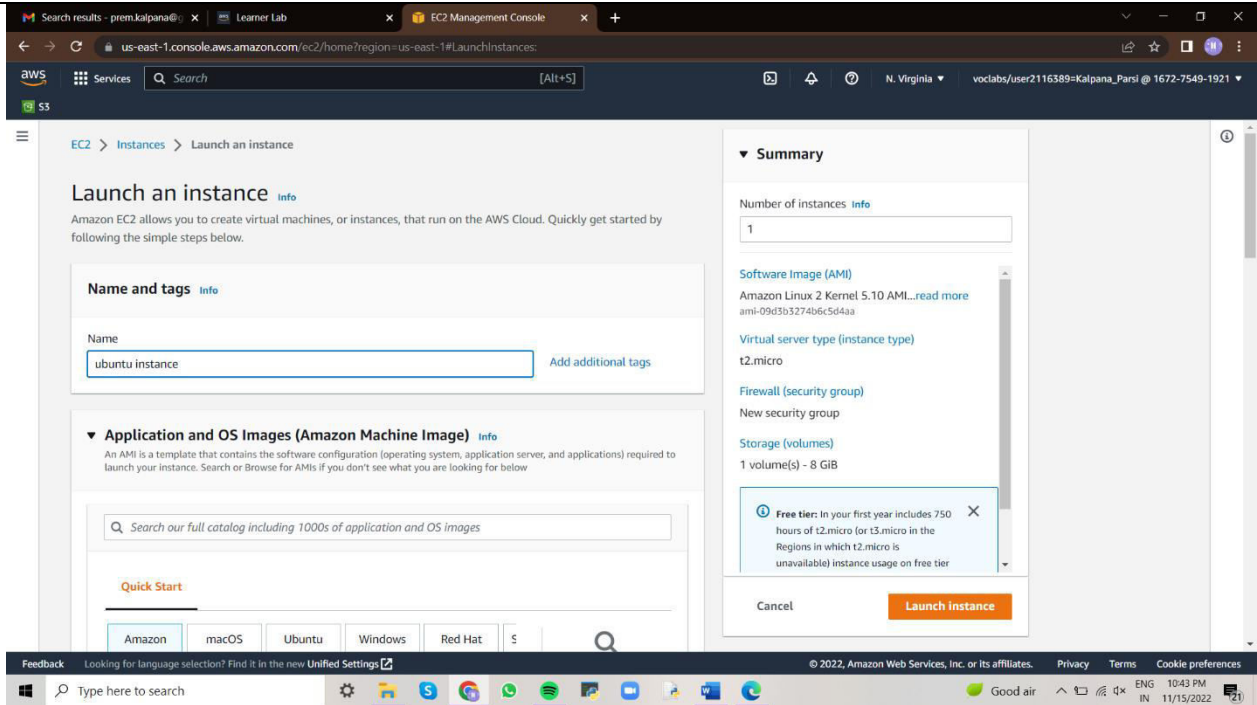
VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

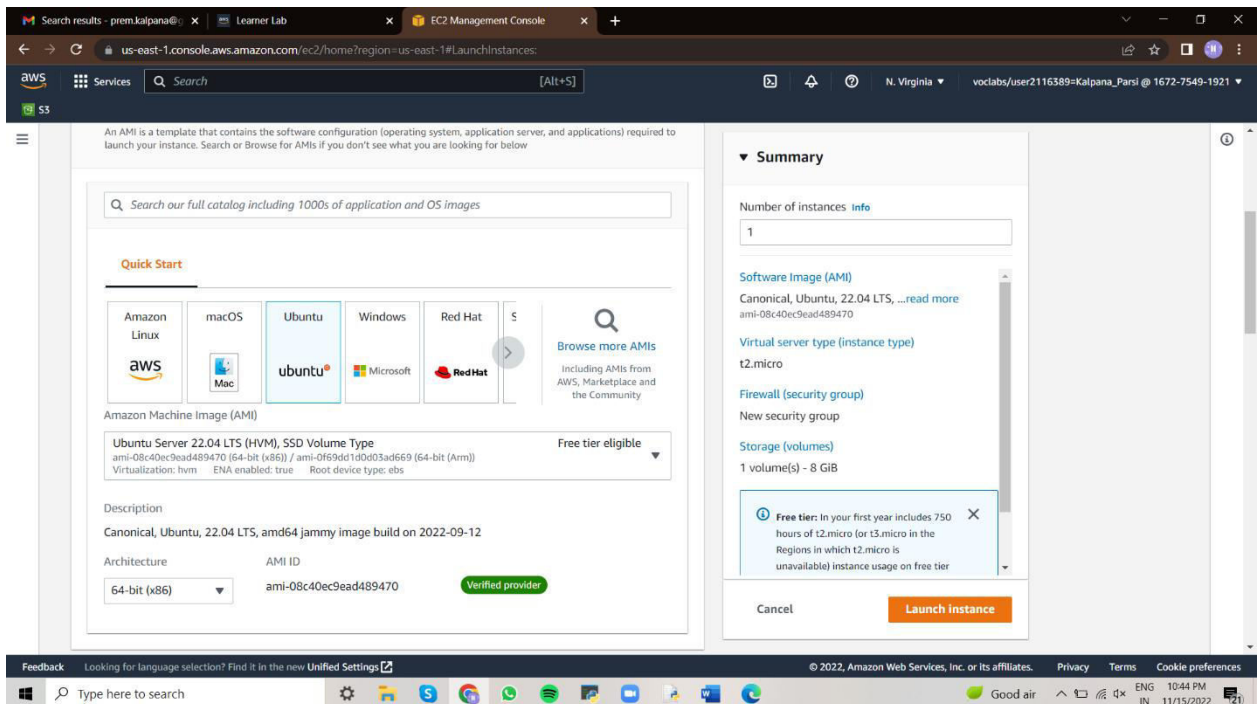
DEPARTMENT OF : Computer Science and Engineering

NAME OF THE LABORATORY : DSCCLAB

Name : _____ Roll No : 1602-19-733-0 Page No: _____



Select Amazon Ubuntu



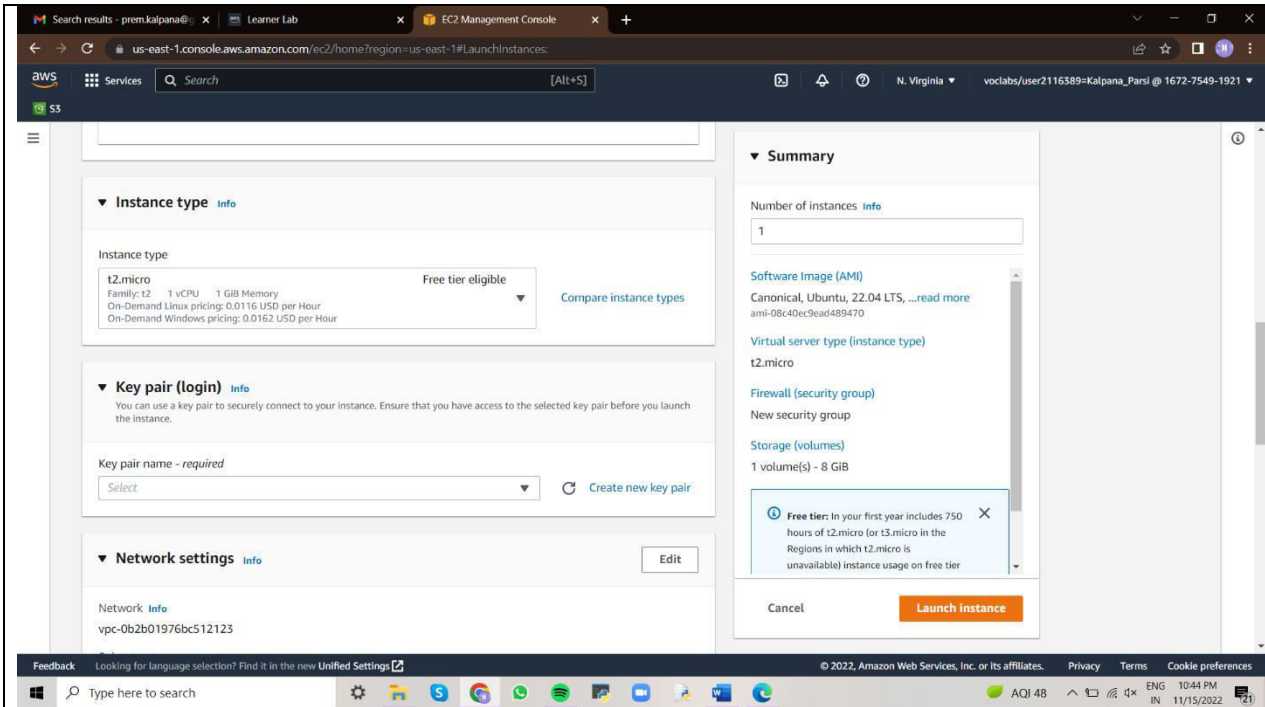
Instance type - t2.micro)

VASAVI COLLEGE OF ENGINEERING

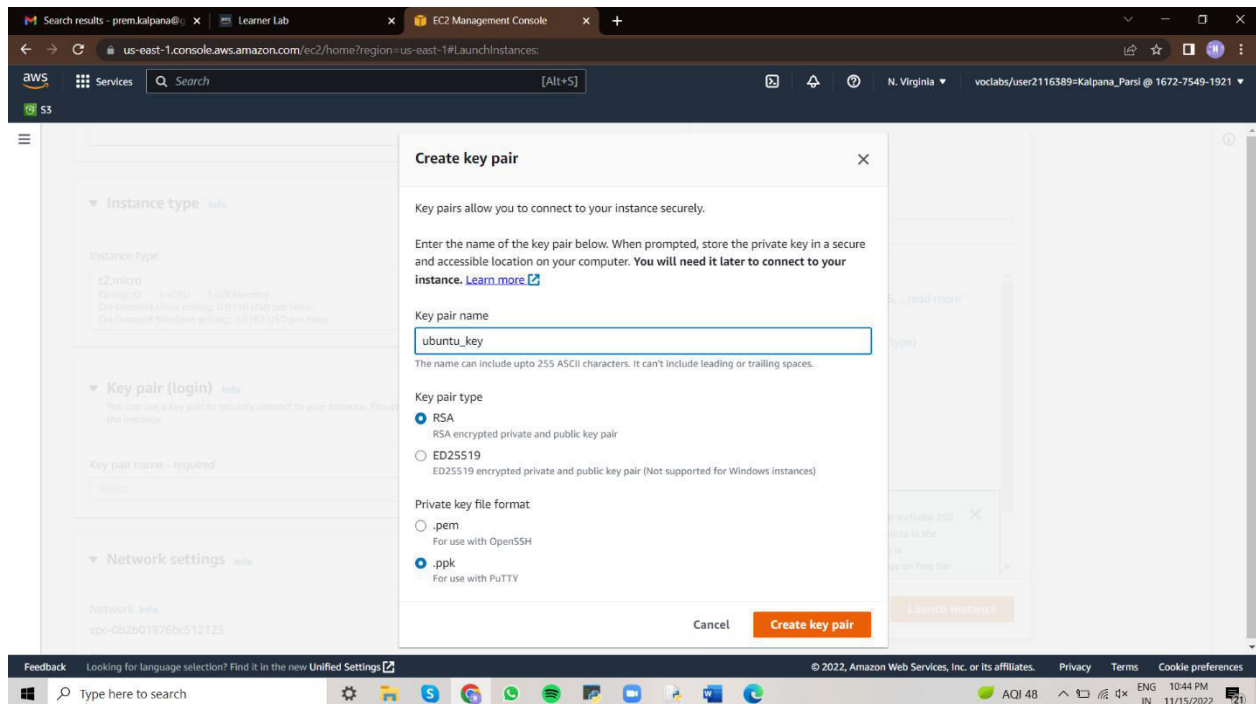
AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCCLAB

Name : _____ Roll No : 1602-19-733-0 Page No: _____



Create new key pair – Save the key pair as .ppk (to work with putty)



Next Add storage

Next configure Security Group – Create security group.

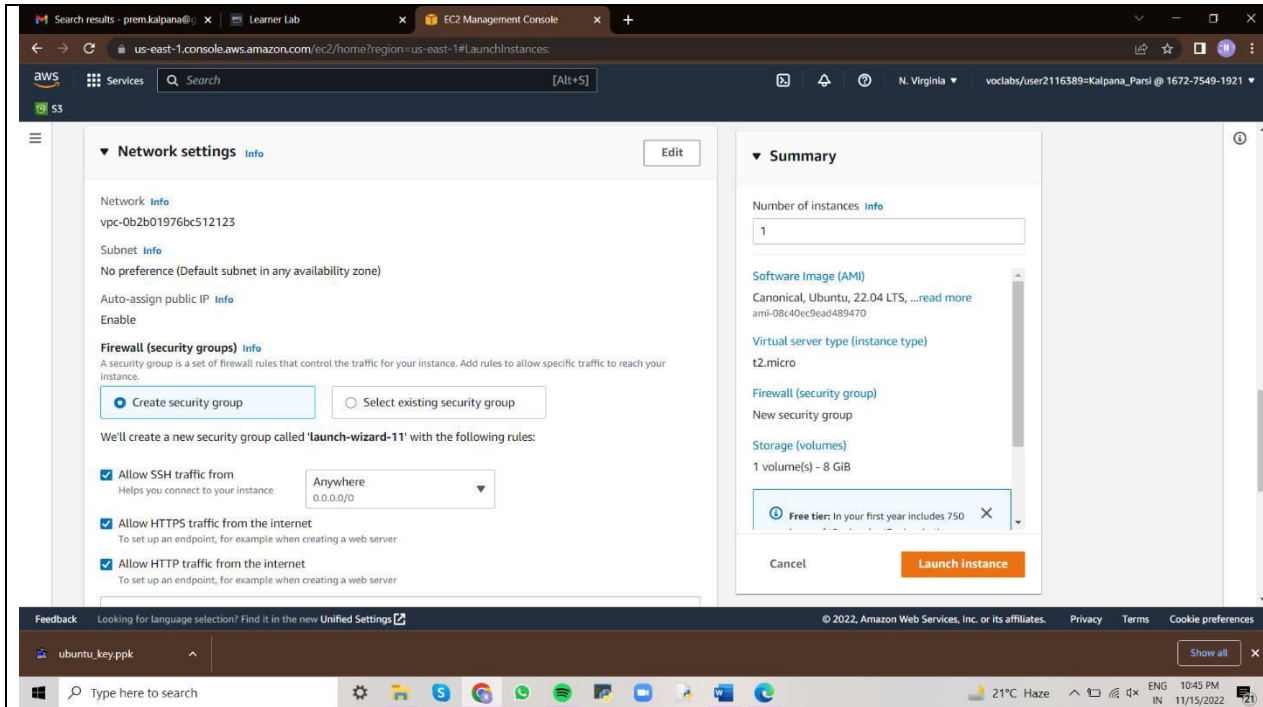
In this step we need to allow http and https requests to access from any group.

VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

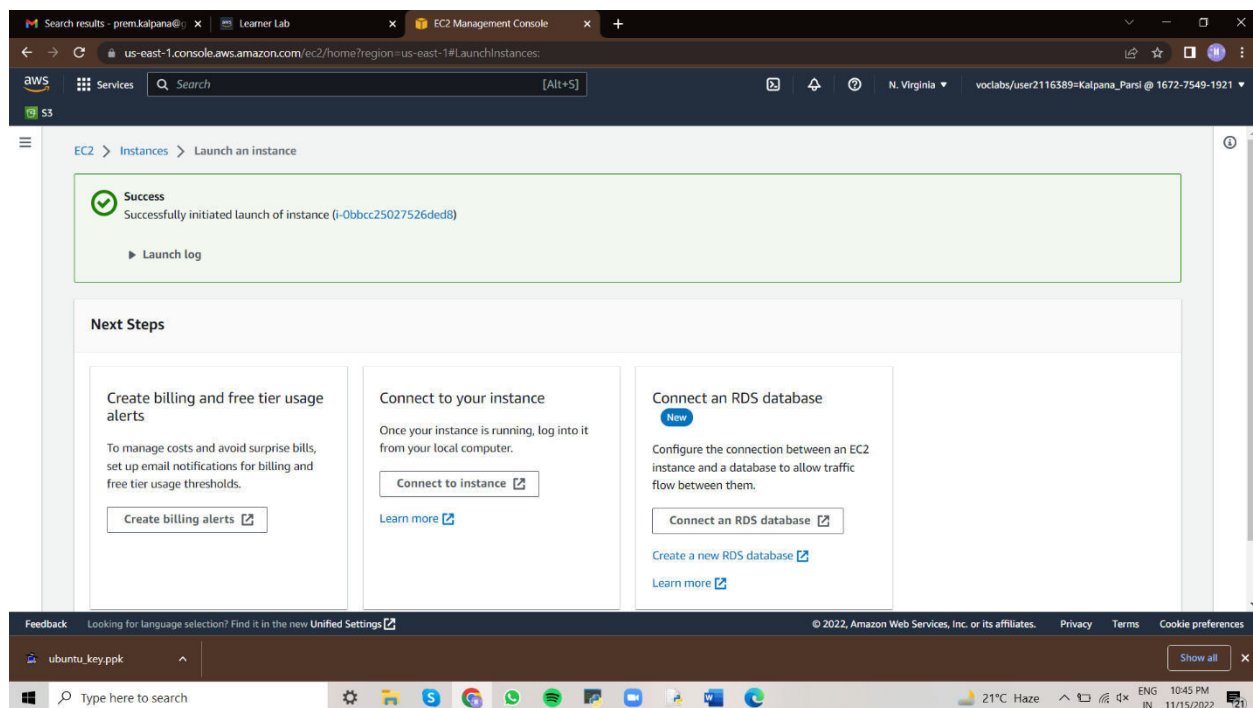
DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCCLAB

Name : _____ Roll No : 1602-19-733-0 Page No: _____



Finally click on Launch instance.

We can see instance is launched successfully.



When the instance state is running , it indicates that your instance was created successfully.

VASAVI COLLEGE OF ENGINEERING

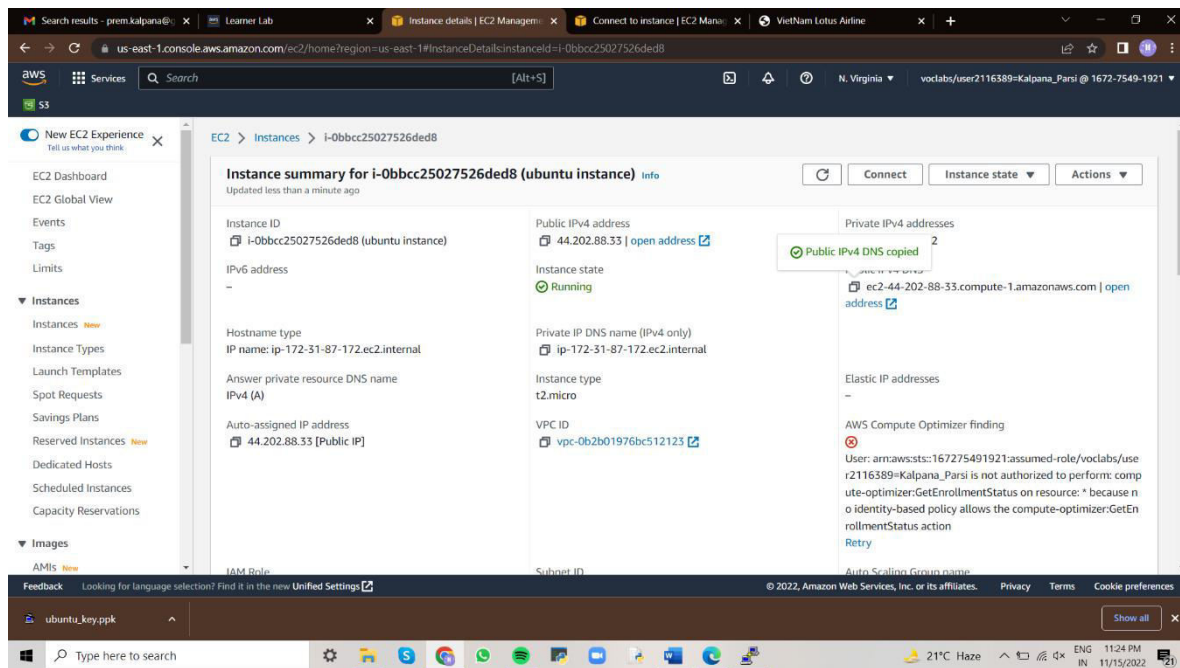
AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering

NAME OF THE LABORATORY : DSCCLAB

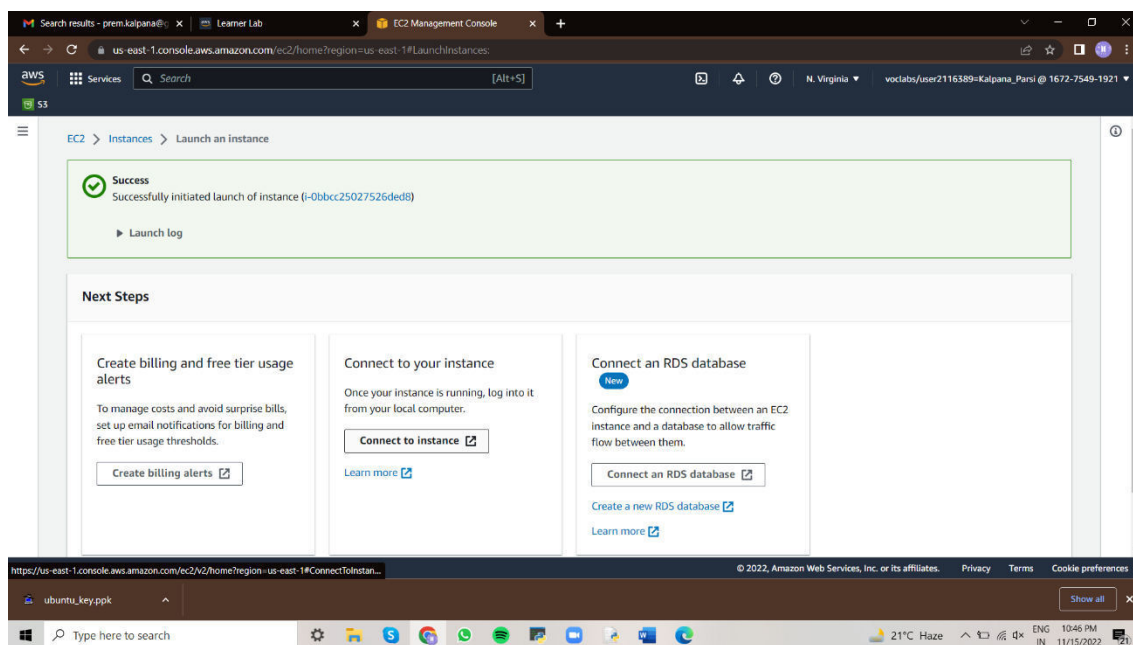
Name : _____ Roll No : 1602-19-733-0 Page No: _____

Copy the public DNS of your Instance. You can access different app running on your instance at a different port.



2. Connect to your Instance:

Click on launch instance then it shows popup window giving details how to connect to your instance.



To open SSH client and If we are in windows platform we need to launch the instance with the help of putty soft.

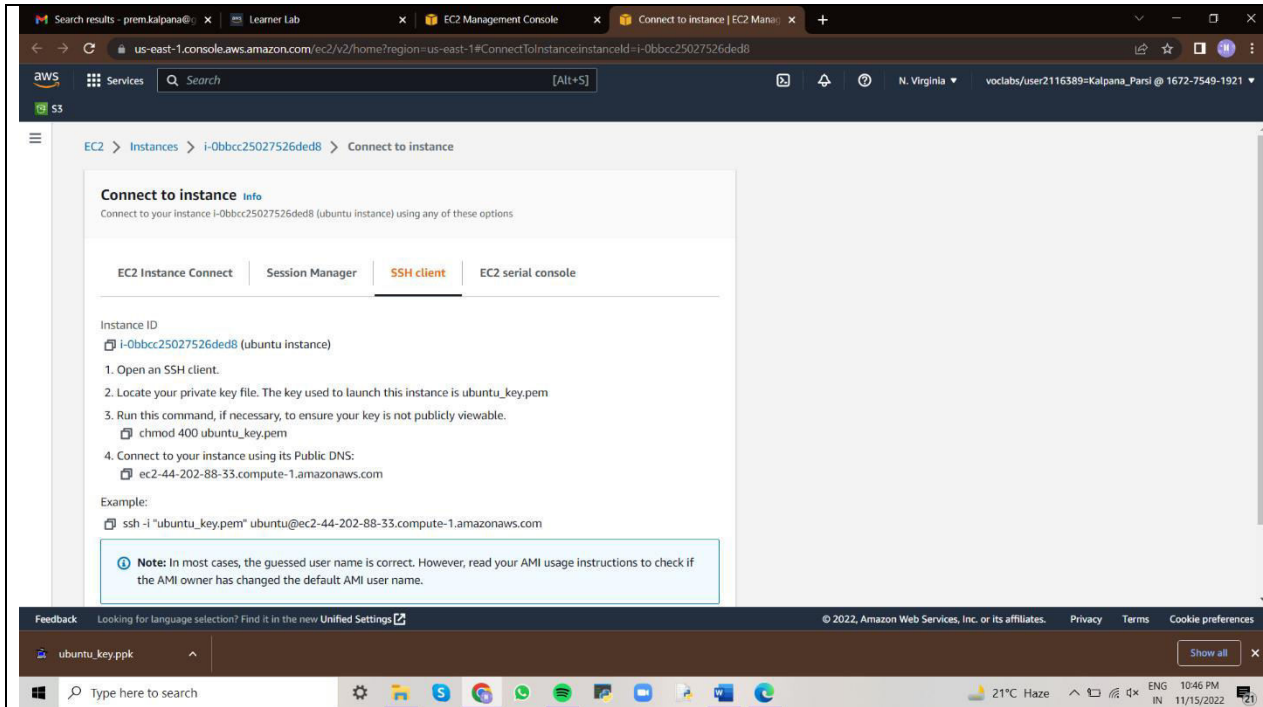
VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

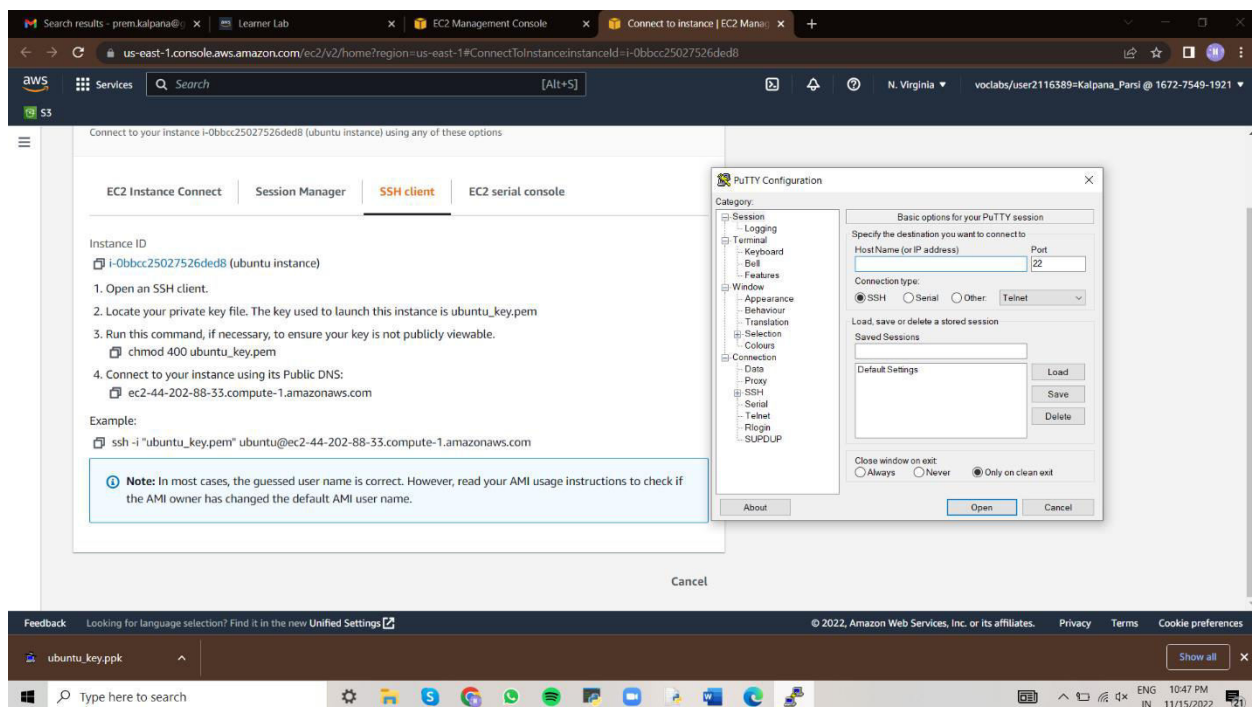
DEPARTMENT OF : Computer Science and Engineering

NAME OF THE LABORATORY : DSCCLAB

Name : _____ Roll No : 1602-19-733-0 Page No: _____



Open Putty



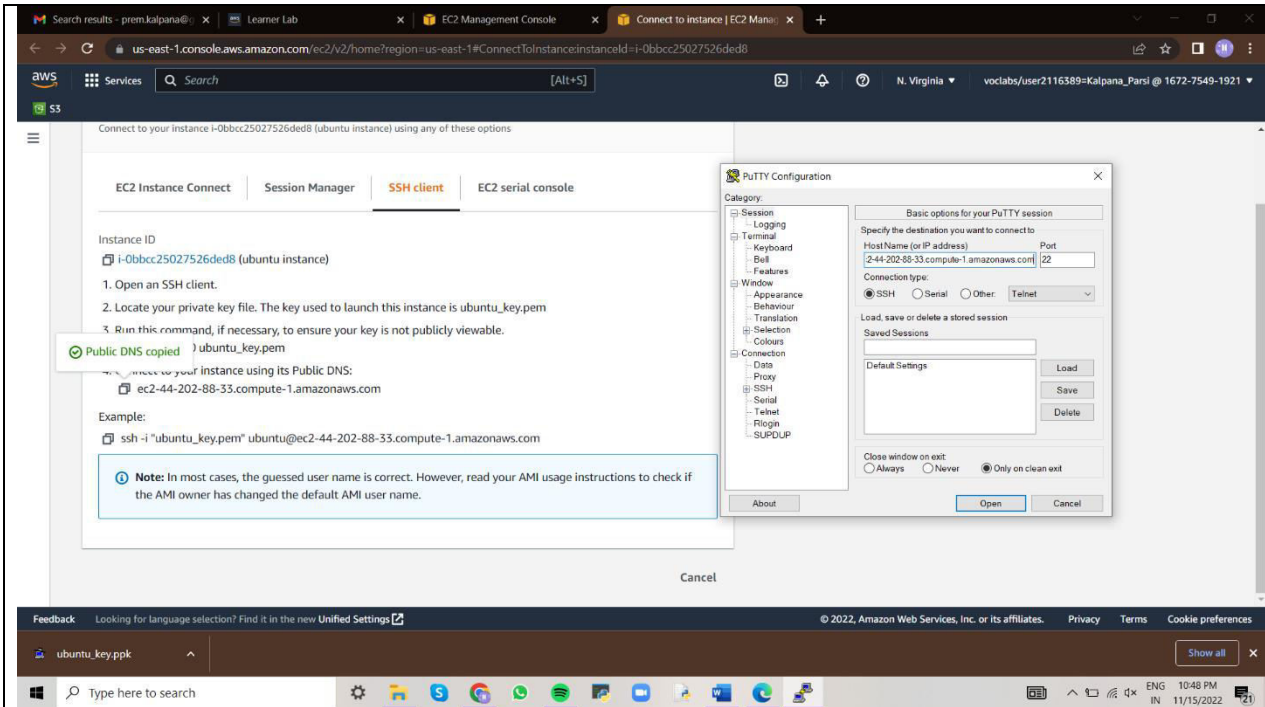
Enter the Public DNS of your Instance in Host Name(IP address)

VASAVI COLLEGE OF ENGINEERING

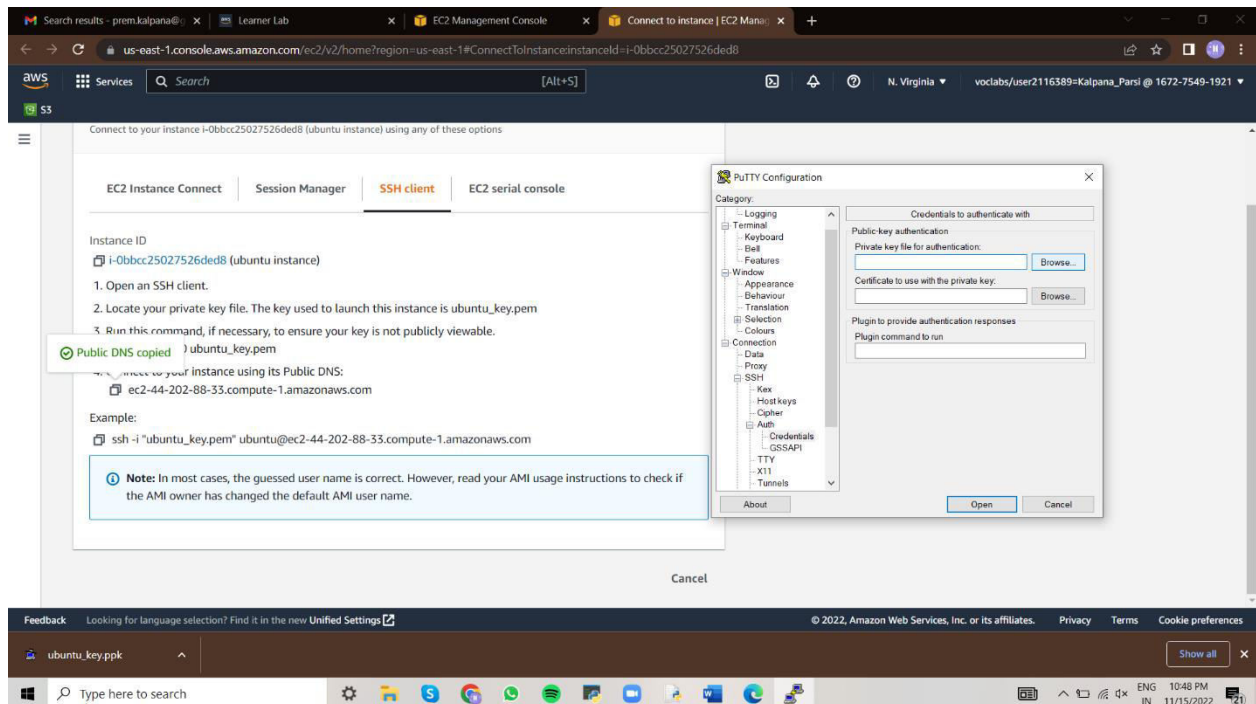
AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCCLAB

Name : _____ Roll No : 1602-19-733-0 Page No: _____



Click on Connection – SSH – Auth – Credentials –



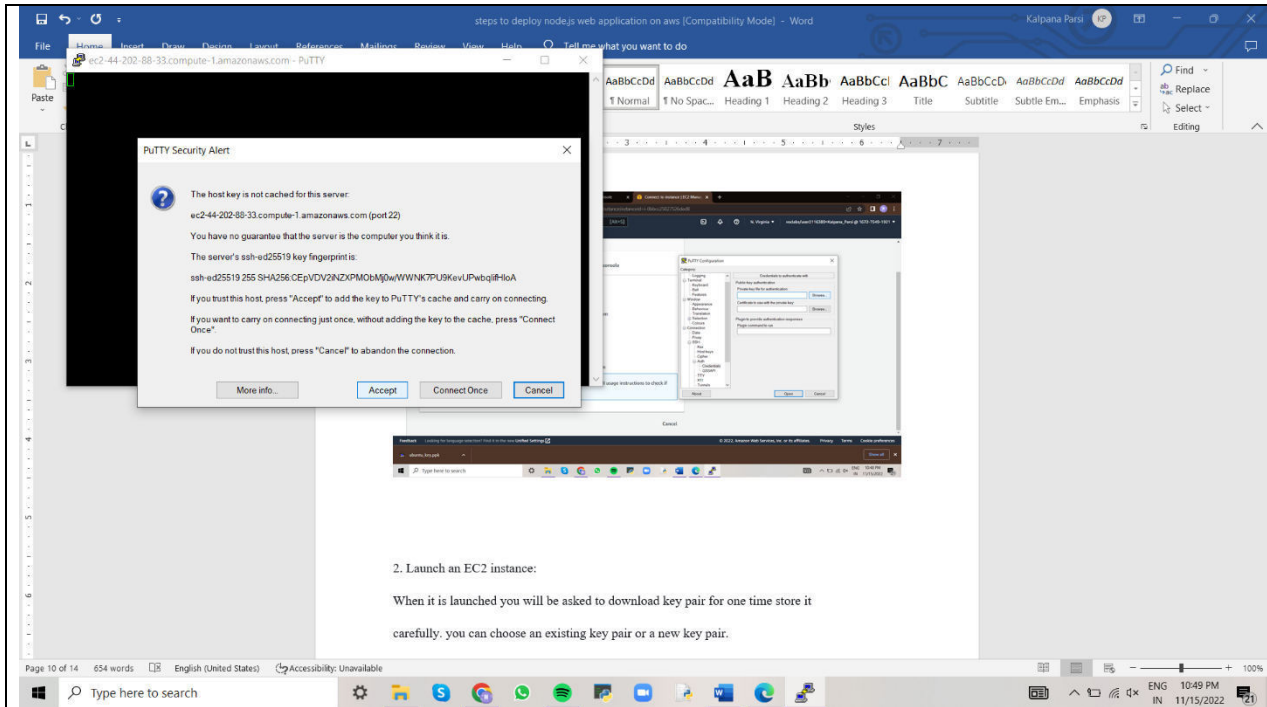
Private key for Authentication - Browse - select the .ppk which was downloaded when EC2 instance is created

VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

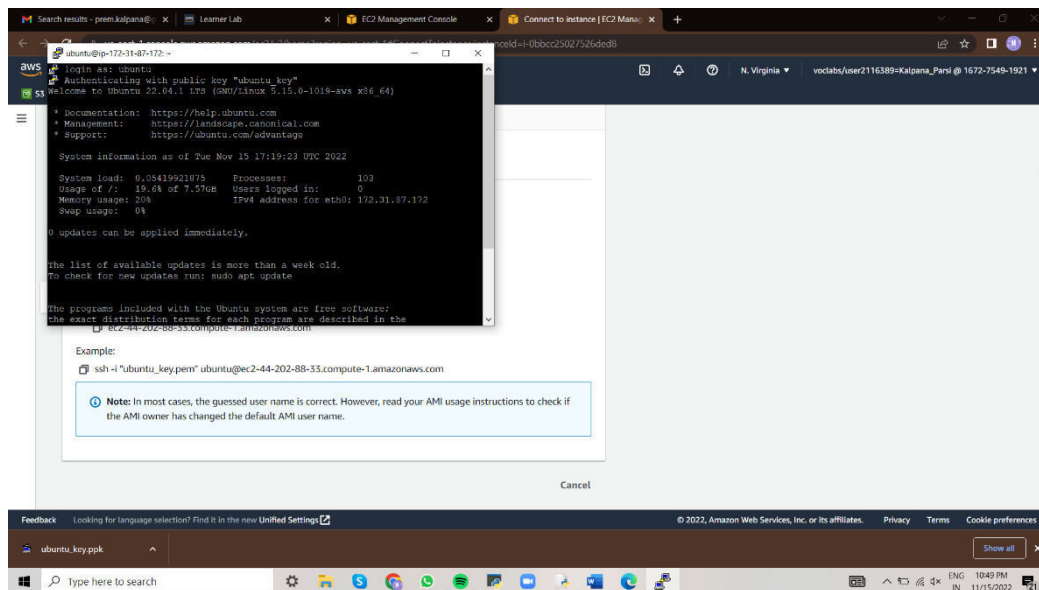
DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCCLAB

Name : _____ Roll No : 1602-19-733-0 Page No: _____



Once entered, it will ask you to confirm, click on Accept

Once it is opened login as ubuntu

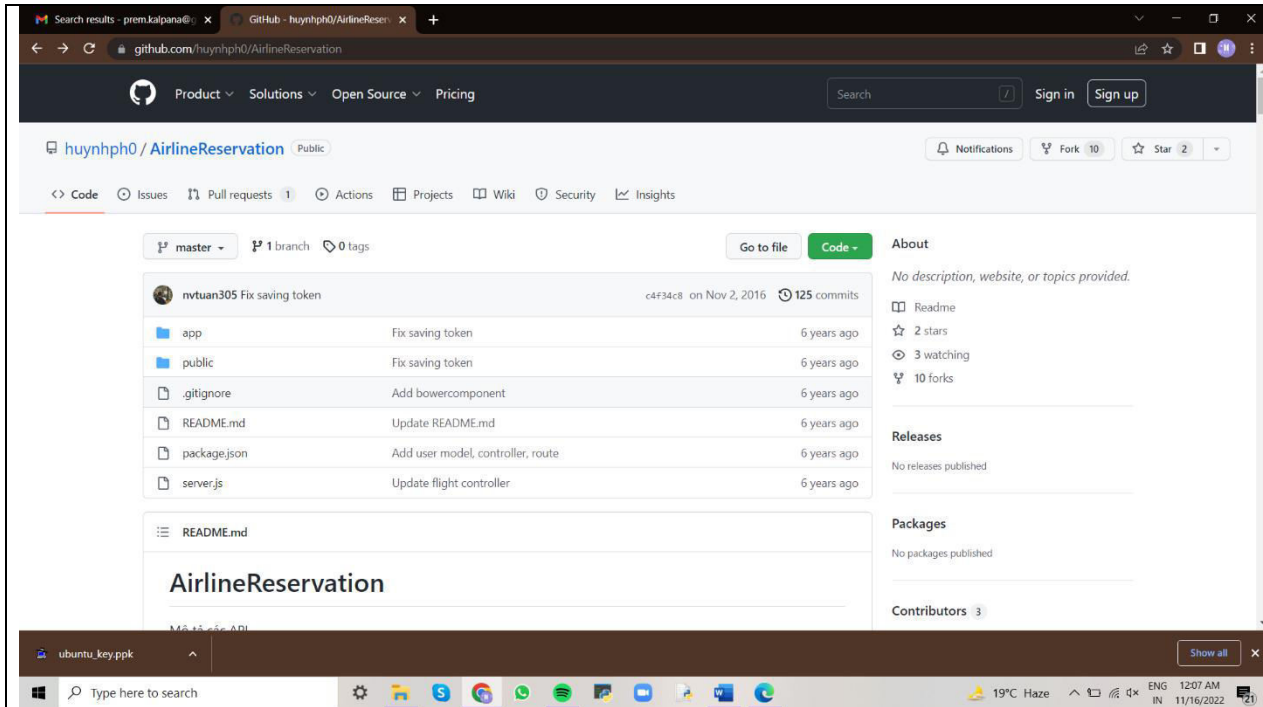


VASAVI COLLEGE OF ENGINEERING

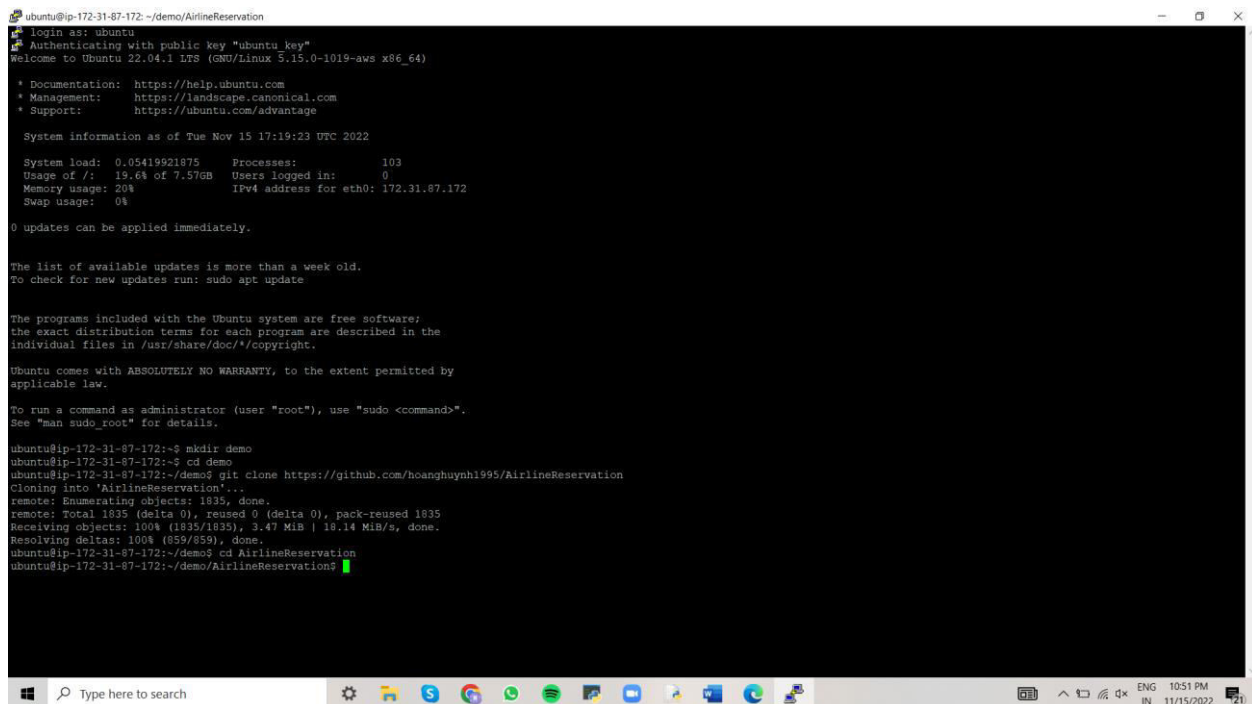
AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCCLAB

Name : _____ Roll No : 1602-19-733-0 Page No: _____



cd AirlineReservation



sudo apt-get update //to download package information from all configured sources

sudo apt-get install npm

//to install Node.js on ubuntu, we must first install npm (node package manager)

VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCCLAB
Name : _____ Roll No : 1602-19-733-0 Page No: ____

select Yes

Ok

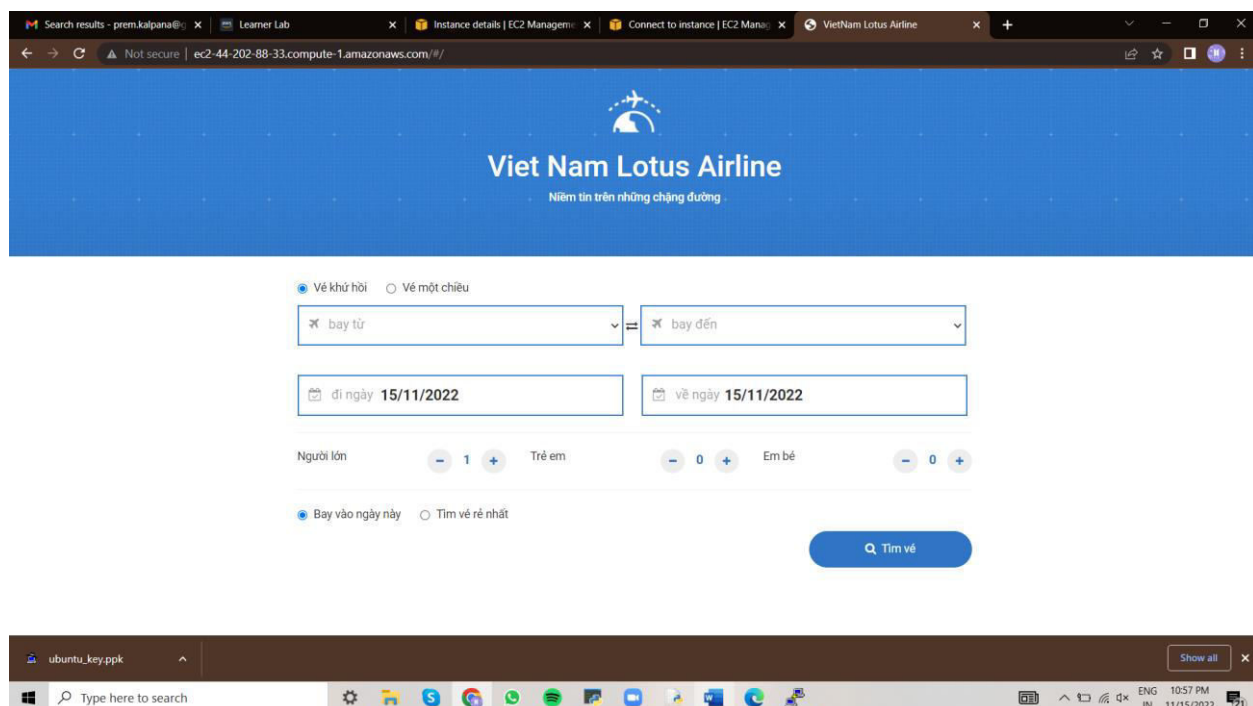
npm install

sudo apt-get install nodejs //to install Node.js on ubuntu

open server.js file using vi editor and change the port no to 80, and save file and exit

sudo node server.js

Copy public DNS of your instance in new tab and view the deployed web application.



VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering

NAME OF THE LABORATORY : DSCCLAB

Name : _____ Roll No : 1602-19-733-0 Page No: _____

LAB PROGRAMS

Implement a distributed application on Hadoop framework.

1. Prerequisites

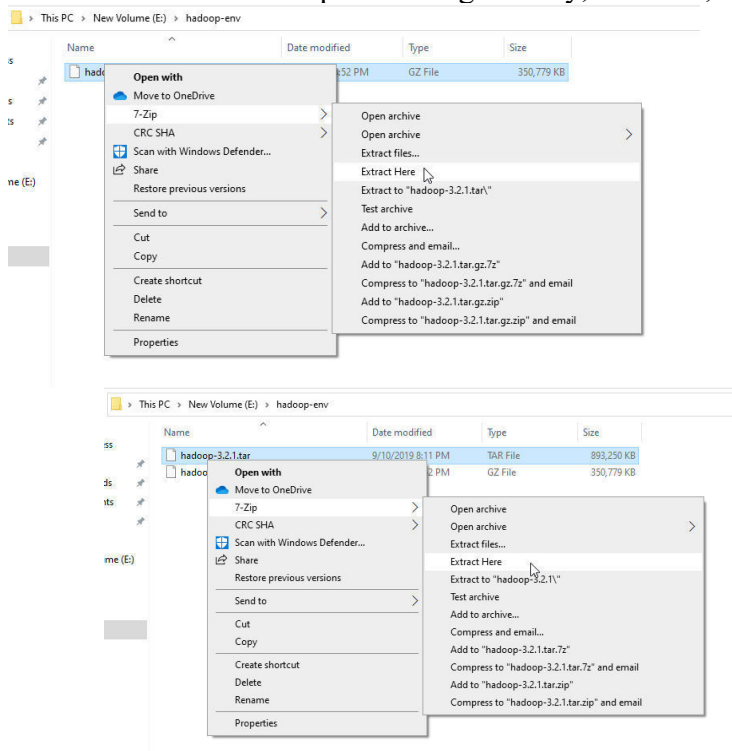
1. Java 8 runtime environment (JRE): Hadoop 3 requires a Java 8 installation. I prefer using the offline installer.
2. Java 8 development Kit (JDK)
3. To unzip downloaded Hadoop binaries, we should install 7zip.

2. Download Hadoop binaries

The first step is to download Hadoop binaries from the official website. The binary package size is about 342 MB.



After finishing the file download, we should unpack the package using 7zip into two steps. First, we should extract the `hadoop-3.2.1.tar.gz` library, and then, we should unpack the extracted tar file:

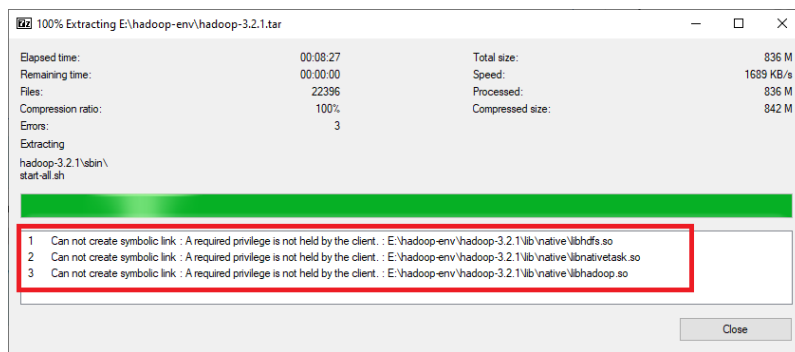


VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCCLAB
Name : _____ Roll No : 1602-19-733-0 Page No: _____

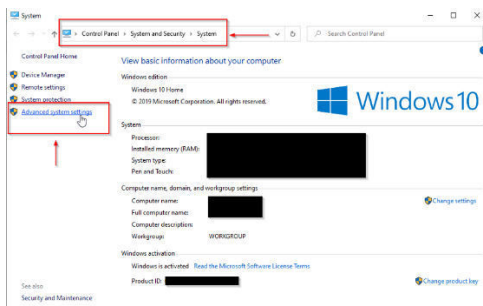
The tar file extraction may take some minutes to finish. In the end, you may see some warnings about symbolic link creation. Just ignore these warnings since they are not related to windows.



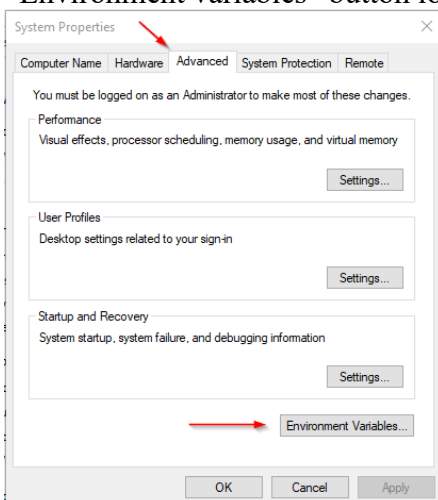
3. Setting up environment variables

After installing Hadoop and its prerequisites, we should configure the environment variables to define Hadoop and Java default paths.

To edit environment variables, go to Control Panel > System and Security > System (or right-click > properties on My Computer icon) and click on the “Advanced system settings” link.



When the “Advanced system settings” dialog appears, go to the “Advanced” tab and click on the “Environment variables” button located on the bottom of the dialog.



VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

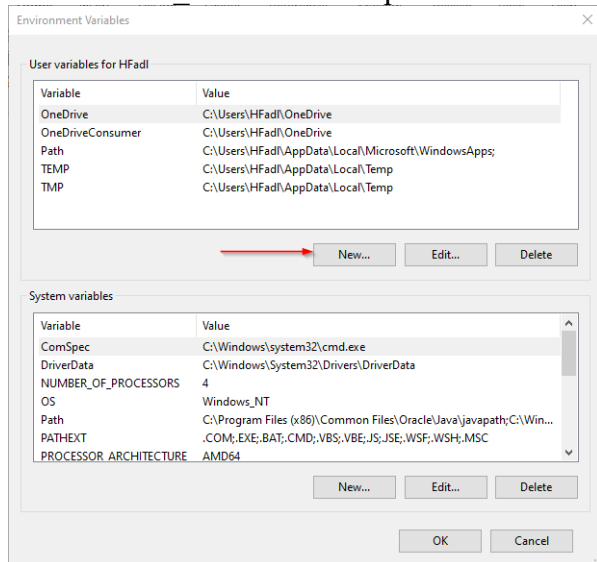
DEPARTMENT OF : Computer Science and Engineering

NAME OF THE LABORATORY : DSCCLAB

Name : _____ Roll No : 1602-19-733-0 Page No: _____

In the “Environment Variables” dialog, press the “New” button to add a new variable. There are two variables to define:

1. JAVA_HOME: JDK installation folder path
2. HADOOP_HOME: Hadoop installation folder path



Now, we should edit the PATH variable to add the Java and Hadoop binaries paths as shown in the following screenshots.




Figure 10 — Editing the PATH variable




Figure 11 — Editing PATH variable

VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering

NAME OF THE LABORATORY : DSCCLAB

Name : _____ Roll No : 1602-19-733-0 Page No: _____




Figure 12— Adding new paths to the PATH variable

4. Configuring Hadoop cluster

There are four files we should alter to configure Hadoop cluster:

1. %HADOOP_HOME%\etc\hadoop\hdfs-site.xml
2. %HADOOP_HOME%\etc\hadoop\core-site.xml
3. %HADOOP_HOME%\etc\hadoop\mapred-site.xml
4. %HADOOP_HOME%\etc\hadoop\yarn-site.xml

4.1. HDFS site configuration

As we know, Hadoop is built using a master-slave paradigm. Before altering the HDFS configuration file, we should create a directory to store all master node (name node) data and another one to store data (data node). In this example, we created the following directories:

- E:\hadoop-env\hadoop-3.2.1\data\dfs\namenode
- E:\hadoop-env\hadoop-3.2.1\data\dfs\datanode

Now, let's open "hdfs-site.xml" file located in "%HADOOP_HOME%\etc\hadoop" directory, and we should add the following properties within

4.2. Core site configuration

Now, we should configure the name node URL adding the following XML code into the <configuration></configuration> element within "core-site.xml":

VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS

(Affiliated to Osmania University)

Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering

NAME OF THE LABORATORY : DSCCLAB

Name : _____ Roll No : 1602-19-733-0 Page No: _____

```
<property><name>fs.default.name</name><value>hdfs://localhost:9820</value></property>
```

4.3. Map Reduce site configuration

Now, we should add the following XML code into the `<configuration></configuration>` element within “mapred-site.xml”:

```
<property><name>mapreduce.framework.name</name><value>yarn</value><description>MapReduce  
framework name</description></property>
```

4.4. Yarn site configuration

Now, we should add the following XML code into the `<configuration></configuration>` element within “yarn-site.xml”:

```
<property><name>yarn.nodemanager.aux-  
services</name><value>mapreduce_shuffle</value><description>Yarn Node Manager Aux  
Service</description></property>
```

5. Formatting Name node

After finishing the configuration, let's try to format the name node using the following command:

`hdfs namenode -format`

Due to a bug you will receive the following error.

This issue will be solved within the next release. For now, you can fix it temporarily using the following steps (reference):

1. Download `hadoop-hdfs-3.2.1.jar` file from the following link.
2. Rename the file name `hadoop-hdfs-3.2.1.jar` to `hadoop-hdfs-3.2.1.bak` in folder `%HADOOP_HOME%\share\hadoop\hdfs`
3. Copy the downloaded `hadoop-hdfs-3.2.1.jar` to folder `%HADOOP_HOME%\share\hadoop\hdfs`

Now, if we try to re-execute the format command (Run the command prompt or PowerShell as administrator), you need to approve file system format.

```
2020-04-17 22:02:58,423 INFO util.GSet: Computing capacity for map NameNodeRetryCache  
2020-04-17 22:02:58,423 INFO util.GSet: VM type = 64-bit  
2020-04-17 22:02:58,424 INFO util.GSet: 0.0299999999329447746% max memory 889 MB = 273.1 KB  
2020-04-17 22:02:58,425 INFO util.GSet: capacity = 2^15 = 32768 entries  
File-format filesystem in Storage Directory root= E:\hadoop-env\hadoop-3.2.1\data\dfs\namenode; location= null ? (Y or N)  
Y
```

Figure 15 — File system format approval

And the command is executed successfully:

```
2020-04-17 22:14:17,206 INFO namenode.FSImage: Allocated new BlockPoolId: BP-2032026115-192.168.1.105-1587150857190  
2020-04-17 22:14:17,207 INFO common.Storage: Will remove files: []  
2020-04-17 22:14:17,275 INFO common.Storage: Storage directory E:\hadoop-env\hadoop-3.2.1\data\dfs\namenode has been suc  
cessfully formatted.  
2020-04-17 22:14:17,331 INFO namenode.FSImageFormatProtobuf: Saving image file E:\hadoop-env\hadoop-3.2.1\data\dfs\namen  
ode\currentFSImage.cpt_000000000000000000 using no compression  
2020-04-17 22:14:17,331 INFO namenode.FSImageFormatProtobuf: Image file E:\hadoop-env\hadoop-3.2.1\data\dfs\namenode\cur  
rentFSImage.cpt_000000000000000000 of size 400 bytes saved in 0 seconds.  
2020-04-17 22:14:17,355 INFO namenode.NamespaceManager: Going to retain 1 images with txid >= 0  
2020-04-17 22:14:17,580 INFO namenode.FSImage: FSImageSaver clean checkpoint: txid=0 when meet shutdown.  
2020-04-17 22:14:17,580 INFO namenode.NameNode: SHUTDOWN_MSG:  
=====SHUTDOWN_MSG: Shutting down NameNode at  
PS C:\Windows\system32>
```

Figure 16 — Command executed successfully

VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCCLAB

Name : _____ Roll No : 1602-19-733-0 Page No: _____

6. Starting Hadoop services

Now, we will open PowerShell, and navigate to “%HADOOP_HOME%\sbin” directory. Then we will run the following command to start the Hadoop nodes:

`.\start-dfs.cmd`



Figure 17 — Starting Hadoop nodes

Two command prompt windows will open (one for the name node and one for the data node) as follows:

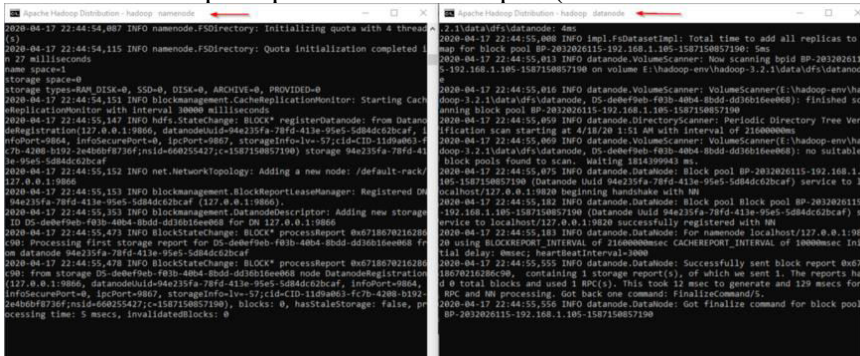


Figure 18 — Hadoop nodes command prompt windows

Next, we must start the Hadoop Yarn service using the following command:

`./start-yarn.cmd`



Figure 19 — Starting Hadoop Yarn services

Two command prompt windows will open (one for the resource manager and one for the node manager) as follows:

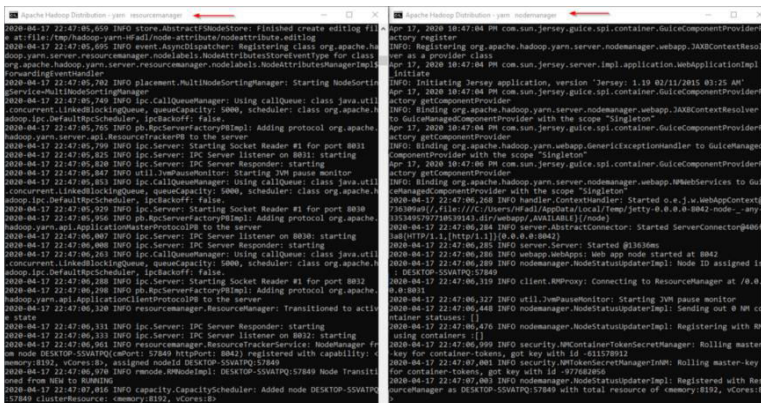


Figure 20— Node manager and Resource manager command prompt windows

VASAVI COLLEGE OF ENGINEERING

AUTONOMOUS
(Affiliated to Osmania University)
Hyderabad- 500 031.

DEPARTMENT OF : Computer Science and Engineering
NAME OF THE LABORATORY : DSCCLAB
Name : _____ Roll No : 1602-19-733-0 Page No: _____

To make sure that all services started successfully, we can run the following command:

Jps

It should display the following services:

```
PS E:\hadoop-env\hadoop-3.2.1\sbin> jps
14560 DataNode
4960 ResourceManager
5936 NameNode
768 NodeManager
14636 Jps
PS E:\hadoop-env\hadoop-3.2.1\sbin>
```

Figure 21 — Executing jps command

7. Hadoop Web UI

There are three web user interfaces to be used:

- Name node web page: <http://localhost:9870/dfshealth.html>

Overview 'localhost:9870' (active)

Started:	Fri Apr 17 22:44:51 +0500 2020
Version:	3.2.1, 653c88a67c22a8229d3808f670c1d07c025842
Compiled:	Tue Sep 10 18:36:00 +0500 2019 by r0hithanmaka from branch-3.2.1
Cluster ID:	CID: 11d5a063-879-4208-8152-2a4b2d8736f
Block Pool ID:	BP-202028115-192-168-1-105-1587150857190

Summary

Security is off.
Safe mode is off.
1 files and directories, 0 blocks (0 replicated blocks, 0 erasure coded block groups) = 1 total filesystem object(s).
Heap Memory used 63.28 MB of 191.5 MB Heap Memory. Max Heap Memory is 889 MB.

- Data node web page: <http://localhost:9864/datanode.html>

DataNode on localhost:9864

Cluster ID:	CID: 11d5a063-879-4208-8152-2a4b2d8736f
Version:	3.2.1, 653c88a67c22a8229d3808f670c1d07c025842

Block Pools

NodeName Address	Block Pool ID	Block State	Last Heartbeat	Last Block Report	Last Block Report Size (Max Size)
localhost:9864	BP-202028115-192-168-1-105-1587150857190	RUNNING	1s	12 minutes	0 B (0 MB)

Volume Information

Directory	StorageType	Capacity Used	Capacity Left	Capacity Reserved	Reserved Space for Replicas	Blocks
-----------	-------------	---------------	---------------	-------------------	-----------------------------	--------

Figure 23 — Data node web page

- Yarn web page: <http://localhost:8088/cluster>

hadoop All Applications

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	V-Cores
0	0	0	0	0	0 B	8 GB	0 B	0

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Re...
1	0	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation
Capacity Scheduler	[memory-mb (unit-M), vcores]	<memory 1024, vCores 1>	<memory 8192, vCores 4>

Show 20 entries

ID	User	Name	Application Type	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU V-Cores	Allocated Memory MB	Reserved CPU V-Cores	Reserved Memory MB
----	------	------	------------------	-------	----------------------	-----------	------------	------------	-------	-------------	--------------------	-----------------------	---------------------	----------------------	--------------------

No data available in table

Showing 0 to 0 of 0 entries

10.Installation and deploying a PhP application on a Docker Container

Description:

Create a Machine Image of Ubuntu Bionic 18.04LTS or Xenial 16.04.

To install **Docker CE**, first, you need to remove older versions of **Docker** were called **docker**, **docker.io**, or **docker-engine** from the system using the following command.

```
$ sudo apt-get remove docker docker-engine docker.io containerd runc
```

Next, you need to set up the Docker repository to install and update Docker from the repository using following commands.

1. Update the apt package index

```
$ sudo apt-get update
```

2. Install packages to allow apt to use a repository over HTTPS

```
$ sudo apt-get install \
    apt-transport-https \
    ca-certificates \
    curl \
    gnupg-agent \
    software-properties-common
```

3. Add Docker's official GPG key

```
$ curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -
```

4. Verify that you now have the key with the fingerprint 9DC8 5822 9FC7 DD38 854A E2D8 8D81 803C 0EBF CD88, by searching for the last 8 characters of the fingerprint

```
sudo apt-key fingerprint 0EBFCD88
```

```
pub rsa4096 2017-02-22 [SCEA]
```

```
9DC8 5822 9FC7 DD38 854A E2D8 8D81 803C 0EBF CD88
```

```
uid      [ unknown] Docker Release (CE deb) <docker@docker.com>
sub  rsa4096 2017-02-22 [S]
```

5. Use the following command to set up the stable repository

```
$ sudo add-apt-repository \
    "deb [arch=amd64] https://download.docker.com/linux/ubuntu \
$(lsb_release -cs) \
    stable"
```

The `lsb_release -cs` sub-command below returns the name of your Ubuntu distribution, such as `xenial`. Sometimes, in a distribution like Linux Mint, you might need to change `$(lsb_release -cs)` to your parent Ubuntu distribution. For example, if you are using Linux Mint Tessa, you could use `bionic`. Docker does not offer any guarantees on untested and unsupported Ubuntu distributions.

6. Update the apt package index and install the latest version of **Docker CE** using following commands.

```
$ sudo apt-get update
```

7. Install the latest version of Docker Engine - Community and `containerd`, or go to the next step 8 to install a specific version

```
$ sudo apt-get install docker-ce docker-ce-cli containerd.io
```

8. To install a specific version of Docker Engine - Community, list the available versions in the repo, then select and install: List the versions available in your repo:

```
$ apt-cache madison docker-ce
```

9. Install a specific version using the version string from the second column, for example, `5:18.09.1~3-0~ubuntu-xenial`

```
$ sudo apt-get install docker-ce=<VERSION_STRING> docker-ce-
cli=<VERSION_STRING> containerd.io
```

10. After successfully installing the **Docker CE** package, the service should be auto-started and auto-enabled to start at system boot, you can check its status using the following command.

```
$ sudo systemctl status docker
```

11. Press CTRL C to exit

12. Verify that Docker Engine - Community is installed correctly by running the hello-world

image

```
$ sudo docker run hello-world
```

13. This command downloads a test image and runs it in a container. When the container runs, it prints the below informational message

Unable to find image 'hello-world:latest' locally

latest: Pulling from library/hello-world

1b930d010525: Pull complete

Digest:

sha256:c3b4ada4687bbaa170745b3e4dd8ac3f194ca95b2d0518b417fb47e5879d9
b5f

Status: Downloaded newer image for hello-world:latest

Hello from Docker!

This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:

1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
(amd64)
3. The Docker daemon created a new container from that image which runs the executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it to your terminal.

To try something more ambitious, you can run an Ubuntu container with:

```
$ docker run -it ubuntu bash
```

Share images, automate workflows, and more with a free Docker ID:

<https://hub.docker.com/>

For more examples and ideas, visit:

<https://docs.docker.com/get-started/>

Dockerizing a Node.js web application

14. Create a new folder namely nodejsapp

15. Make a package.json file as follows

```
{
  "name": "docker_web_app",
  "version": "1.0.0",
  "description": "Node.js on Docker",
  "author": "Sashi's First Nodejs Application on Container
    <sashi.mamidanna@gmail.com>",
  "main": "server.js",
  "scripts": {
    "start": "node server.js"
  },
  "dependencies": {
    "express": "^4.16.1"
  }
}
```

16. Then create a file server.js to create a program that runs on the node. The idea is to enable the server.js file to run on the container at port no 8081

```
'use strict';
const express = require('express');
```

```
// Constants
```

```
const PORT = 8081;
const HOST = '0.0.0.0';
```

```
// App
```

```
const app = express();
app.get('/', (req, res) => {
  res.send('Hello world\n');
});
```

```
app.listen(PORT, HOST);
```

```
console.log(`Running on http://${HOST}:${PORT}`);
```

17. Create a dockerfile now namely dockerfile in the same directory

```
$sudo nano dockerfile
```

18. Copy the source code into the dockerfile

```
FROM node:10
```

```
# Create app directory
```

```
WORKDIR /app
```

```
COPY . /app
```

```
RUN npm install
```

```
COPY . .
```

```
EXPOSE 8082
```

```
CMD [ "node", "server.js" ]
```

19. Now build the docker image with the node application on it

```
$sudo docker build -t nodejsapp .
```

20. Run the application by executing run command on docker

```
$sudo docker run -p 8082:8081 nodejsapp
```

21. The container engine will run the command node server.js that was initialized through the dockerfile. Now the server.js is listening to incoming requests on <http://localhost:8081> on the host operating system. But the application is running on port number 8082 on the docker engine.

22. Open a new ssh connection on the same VM and run the command to send an outgoing request to the application running on docker

```
$sudo curl http://localhost:8082
```

```
Hello World
```

23. This response is a result of the application running on node, devoted on the docker container, that's running on Docker engine available on the Ubuntu OS.
24. Run the bow command to check if the docker image is present in the list of images on Docker C

```
$sudo docker ps
```
25. To stop the docker container image

```
$sudo docker stop <docker image ID>
```
26. To remove the docker image
27.

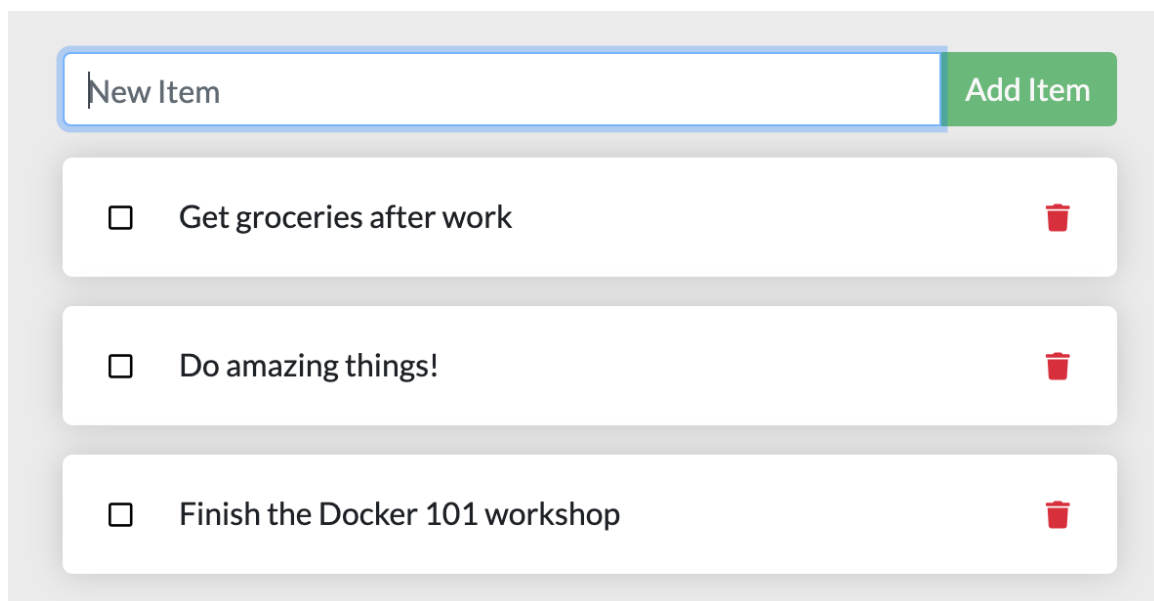
```
$sudo docker rmi <docker image ID>
```


12. Installation and deploying a node.js application on a Docker Container

1. Install Docker desktop and Git client

For the rest of this tutorial, we will be working with a simple todo list manager that is running in Node.js. If you're not familiar with Node.js, don't worry! No real JavaScript experience is needed!

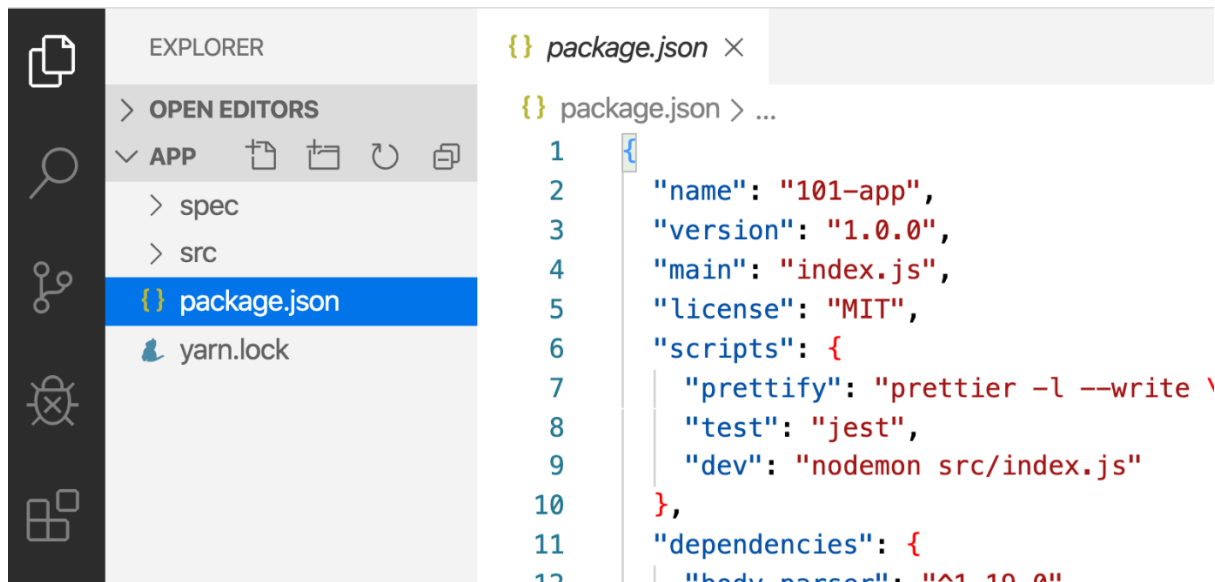
At this point, your development team is quite small and you're simply building an app to prove out your MVP (minimum viable product). You want to show how it works and what it's capable of doing without needing to think about how it will work for a large team, multiple developers, etc.



Getting our App¶

Before we can run the application, we need to get the application source code onto our machine. For real projects, you will typically clone the repo. But, for this tutorial, we have created a ZIP file containing the application.

1. [Download the ZIP](#). Open the ZIP file and make sure you extract the contents.
2. Once extracted, use your favorite code editor to open the project. If you're in need of an editor, you can use [Visual Studio Code](#). You should see the `package.json` and two subdirectories (`src` and `spec`).



Building the App's Container Image

In order to build the application, we need to use a `Dockerfile`. A `Dockerfile` is simply a text-based script of instructions that is used to create a container image. If you've created `Dockerfiles` before, you might see a few flaws in the `Dockerfile` below. But, don't worry! We'll go over them.

1. Create a file named `Dockerfile` in the same folder as the file `package.json` with the following contents.

```
2. FROM node:18-alpine
3. WORKDIR /app
4. COPY . .
5. RUN yarn install --production
6. CMD ["node", "src/index.js"]
```

Please check that the file `Dockerfile` has no file extension like `.txt`. Some editors may append this file extension automatically and this would result in an error in the next step.

7. If you haven't already done so, open a terminal and go to the `app` directory with the `Dockerfile`. Now build the container image using the `docker build` command.
8. `docker build -t getting-started .`

This command used the `Dockerfile` to build a new container image. You might have noticed that a lot of "layers" were downloaded. This is because we instructed the builder that we wanted to start from the `node:18-alpine` image. But, since we didn't have that on our machine, that image needed to be downloaded.

After the image was downloaded, we copied in our application and used `yarn` to install our application's dependencies. The `CMD` directive specifies the default command to run when starting a container from this image.

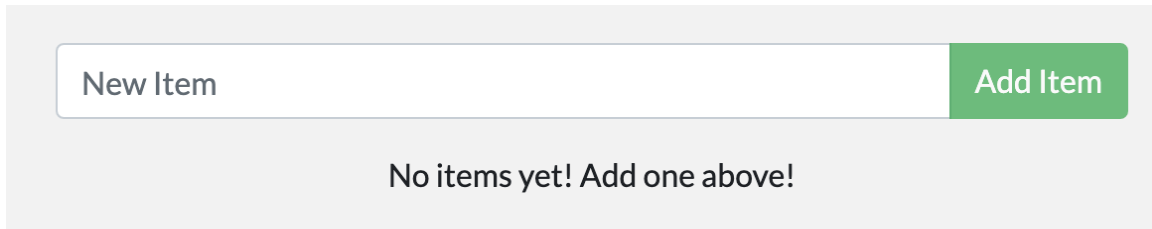
Finally, the `-t` flag tags our image. Think of this simply as a human-readable name for the final image. Since we named the image `getting-started`, we can refer to that image when we run a container.

The `.` at the end of the `docker build` command tells that Docker should look for the `Dockerfile` in the current directory.

Starting an App Container¶

Now that we have an image, let's run the application! To do so, we will use the `docker run` command (remember that from earlier?).

1. Start your container using the `docker run` command and specify the name of the image we just created:
2. `docker run -dp 3000:3000 getting-started`
Remember the `-d` and `-p` flags? We're running the new container in "detached" mode (in the background) and creating a mapping between the host's port 3000 to the container's port 3000. Without the port mapping, we wouldn't be able to access the application.
3. After a few seconds, open your web browser to <http://localhost:3000>. You should see our app!

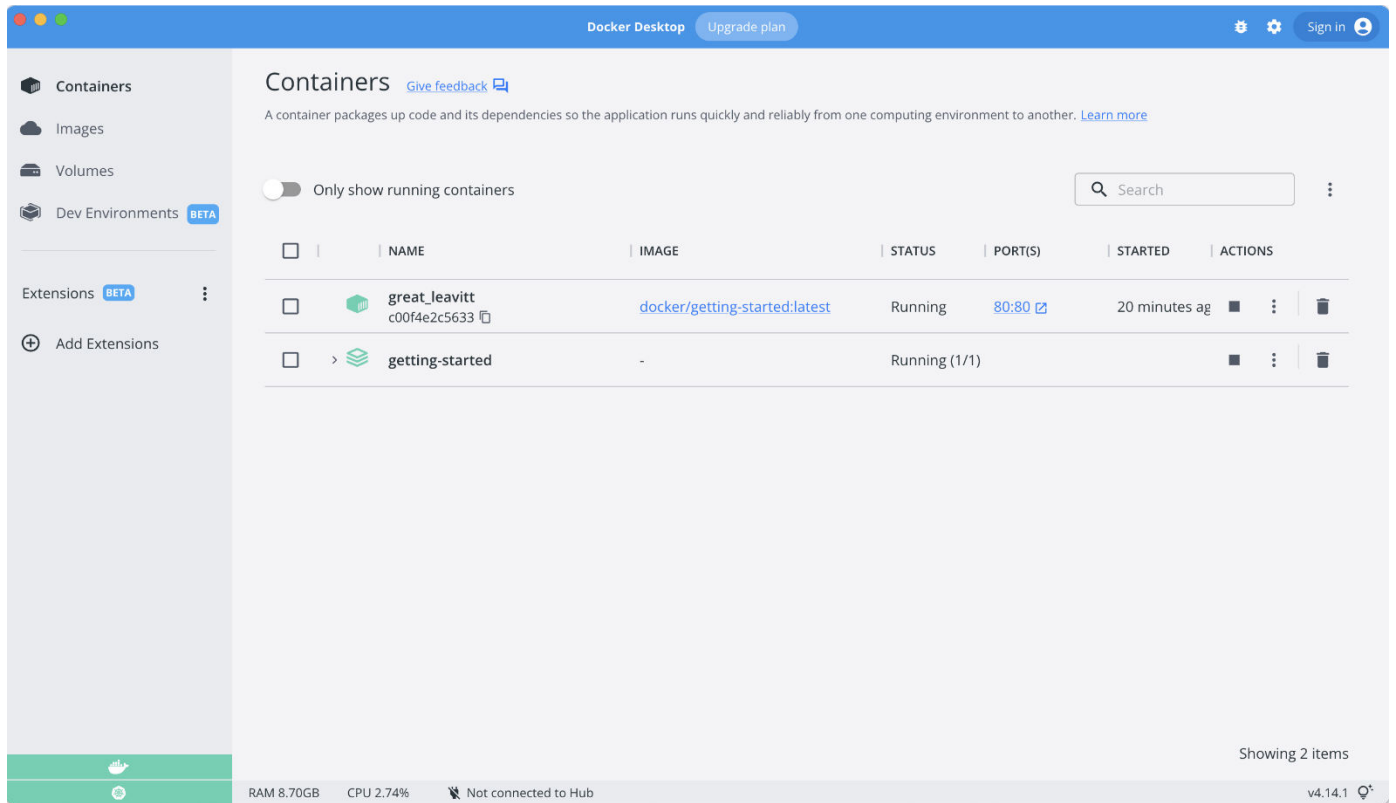


The screenshot shows a web application interface. At the top, there is a light gray box containing a white input field with the placeholder text "New Item" and a green button labeled "Add Item". Below this box, centered on the page, is the text "No items yet! Add one above!" in a bold, black font.

4. Go ahead and add an item or two and see that it works as you expect. You can mark items as complete and remove items. Your frontend is successfully storing items in the backend! Pretty quick and easy, huh?

At this point, you should have a running todo list manager with a few items, all built by you! Now, let's make a few changes and learn about managing our containers.

If you take a quick look at the Docker Dashboard, you should see your two containers running now (this tutorial and your freshly launched app container)!



Recap¶

In this short section, we learned the very basics about building a container image and created a Dockerfile to do so. Once we built an image, we started the container and saw the running app!